

```

0 LightGlue/benchmark.py
1 # Benchmark script for LightGlue on real images
2 import argparse
3 import time
4 from collections import defaultdict
5 from pathlib import Path
6
7 import matplotlib.pyplot as plt
8 import numpy as np
9 import torch
10 import torch._dynamo
11
12 from lightglue import LightGlue, SuperPoint
13 from lightglue.utils import load_image
14
15 torch.set_grad_enabled(False)
16
17
18 def measure(matcher, data, device="cuda", r=100):
19     timings = np.zeros((r, 1))
20     if device.type == "cuda":
21         starter = torch.cuda.Event(enable_timing=True)
22         ender = torch.cuda.Event(enable_timing=True)
23     # warmup
24     for _ in range(10):
25         _ = matcher(data)
26     # measurements
27     with torch.no_grad():
28         for rep in range(r):
29             if device.type == "cuda":
30                 starter.record()
31                 _ = matcher(data)
32                 ender.record()
33                 # sync gpu
34                 torch.cuda.synchronize()
35                 curr_time = starter.elapsed_time(ender)
36             else:
37                 start = time.perf_counter()
38                 _ = matcher(data)
39                 curr_time = (time.perf_counter() - start) * 1e3
40             timings[rep] = curr_time
41     mean_syn = np.sum(timings) / r
42     std_syn = np.std(timings)
43     return {"mean": mean_syn, "std": std_syn}
44
45
46 def print_as_table(d, title, cnames):
47     print()
48     header = f"{title:30} " + " ".join([f"{x:>7}" for x in cnames])
49     print(header)
50     print("-" * len(header))
51     for k, l in d.items():
52         print(f"{k:30}", " ".join([f"{x:>7.1f}" for x in l]))
53
54
55 if __name__ == "__main__":
56     parser = argparse.ArgumentParser(description="Benchmark script for LightGlue")
57     parser.add_argument(
58         "--device",
59         choices=["auto", "cuda", "cpu", "mps"],
60         default="auto",
61         help="device to benchmark on",
62     )
63     parser.add_argument("--compile", action="store_true", help="Compile LightGlue runs")
64     parser.add_argument(
65         "--no_flash", action="store_true", help="disable FlashAttention"
66     )
67     parser.add_argument(
68         "--no_prune_thresholds",
69         action="store_true",
70         help="disable pruning thresholds (i.e. always do pruning)",
71     )
72     parser.add_argument(
73         "--add_superglue",
74         action="store_true",
75         help="add SuperGlue to the benchmark (requires hloc)",
76     )
77     parser.add_argument(
78         "--measure", default="time", choices=["time", "log-time", "throughput"]
79     )
80     parser.add_argument(
81         "--repeat", "--r", type=int, default=100, help="repetitions of measurements"
82     )

```

```

83 parser.add_argument(
84     "--num_keypoints",
85     nargs="+",
86     type=int,
87     default=[256, 512, 1024, 2048, 4096],
88     help="number of keypoints (list separated by spaces)",
89 )
90 parser.add_argument(
91     "--matmul_precision", default="highest", choices=["highest", "high", "medium"]
92 )
93 parser.add_argument(
94     "--save", default=None, type=str, help="path where figure should be saved"
95 )
96 args = parser.parse_intermixed_args()
97
98 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
99 if args.device != "auto":
100     device = torch.device(args.device)
101
102 print("Running benchmark on device:", device)
103
104 images = Path("assets")
105 inputs = {
106     "easy": (
107         load_image(images / "DSC_0411.JPG"),
108         load_image(images / "DSC_0410.JPG"),
109     ),
110     "difficult": (
111         load_image(images / "sacre_coeur1.jpg"),
112         load_image(images / "sacre_coeur2.jpg"),
113     ),
114 }
115
116 configs = {
117     "LightGlue-full": {
118         "depth_confidence": -1,
119         "width_confidence": -1,
120     },
121     # 'LG-prune': {
122     #     'width_confidence': -1,
123     # },
124     # 'LG-depth': {
125     #     'depth_confidence': -1,
126     # },
127     "LightGlue-adaptive": {},
128 }
129
130 if args.compile:
131     configs = {**configs, **{k + "-compile": v for k, v in configs.items()}}
132
133 sg_configs = {
134     # 'SuperGlue': {},
135     "SuperGlue-fast": {"sinkhorn_iterations": 5}
136 }
137
138 torch.set_float32_matmul_precision(args.matmul_precision)
139
140 results = {k: defaultdict(list) for k, v in inputs.items()}
141
142 extractor = SuperPoint(max_num_keypoints=None, detection_threshold=-1)
143 extractor = extractor.eval().to(device)
144 figsize = (len(inputs) * 4.5, 4.5)
145 fig, axes = plt.subplots(1, len(inputs), sharey=True, figsize=figsize)
146 axes = axes if len(inputs) > 1 else [axes]
147 fig.canvas.manager.set_window_title(f"LightGlue benchmark ({device.type})")
148
149 for title, ax in zip(inputs.keys(), axes):
150     ax.set_xscale("log", base=2)
151     bases = [2**x for x in range(7, 16)]
152     ax.set_xticks(bases, bases)
153     ax.grid(which="major")
154     if args.measure == "log-time":
155         ax.set_yscale("log")
156         yticks = [10**x for x in range(6)]
157         ax.set_yticks(yticks, yticks)
158         mpos = [10**x * i for x in range(6) for i in range(2, 10)]
159         mlabel = [
160             10**x * i if i in [2, 5] else None
161             for x in range(6)
162             for i in range(2, 10)
163         ]
164         ax.set_yticks(mpos, mlabel, minor=True)
165         ax.grid(which="minor", linewidth=0.2)
166     ax.set_title(title)
167

```

```

168     ax.set_xlabel("# keypoints")
169     if args.measure == "throughput":
170         ax.set_ylabel("Throughput [pairs/s]")
171     else:
172         ax.set_ylabel("Latency [ms]")
173
174 for name, conf in configs.items():
175     print("Run benchmark for:", name)
176     torch.cuda.empty_cache()
177     matcher = LightGlue(features="superpoint", flash=not args.no_flash, **conf)
178     if args.no_prune_thresholds:
179         matcher.pruning_keypoint_thresholds = {
180             k: -1 for k in matcher.pruning_keypoint_thresholds
181         }
182     matcher = matcher.eval().to(device)
183     if name.endswith("compile"):
184         import torch._dynamo
185
186         torch._dynamo.reset() # avoid buffer overflow
187         matcher.compile()
188     for pair_name, ax in zip(inputs.keys(), axes):
189         image0, image1 = [x.to(device) for x in inputs[pair_name]]
190         runtimes = []
191         for num_kpts in args.num_keypoints:
192             extractor.conf.max_num_keypoints = num_kpts
193             feats0 = extractor.extract(image0)
194             feats1 = extractor.extract(image1)
195             runtime = measure(
196                 matcher,
197                 {"image0": feats0, "image1": feats1},
198                 device=device,
199                 r=args.repeat,
200             )["mean"]
201             results[pair_name][name].append(
202                 1000 / runtime if args.measure == "throughput" else runtime
203             )
204         ax.plot(
205             args.num_keypoints, results[pair_name][name], label=name, marker="o"
206         )
207     del matcher, feats0, feats1
208
209 if args.add_superglue:
210     from hloc.matchers.superglue import SuperGlue
211
212     for name, conf in sg_configs.items():
213         print("Run benchmark for:", name)
214         matcher = SuperGlue(conf)
215         matcher = matcher.eval().to(device)
216         for pair_name, ax in zip(inputs.keys(), axes):
217             image0, image1 = [x.to(device) for x in inputs[pair_name]]
218             runtimes = []
219             for num_kpts in args.num_keypoints:
220                 extractor.conf.max_num_keypoints = num_kpts
221                 feats0 = extractor.extract(image0)
222                 feats1 = extractor.extract(image1)
223                 data = {
224                     "image0": image0[None],
225                     "image1": image1[None],
226                     **{k + "0": v for k, v in feats0.items()},
227                     **{k + "1": v for k, v in feats1.items()},
228                 }
229                 data["scores0"] = data["keypoint_scores0"]
230                 data["scores1"] = data["keypoint_scores1"]
231                 data["descriptors0"] = (
232                     data["descriptors0"].transpose(-1, -2).contiguous()
233                 )
234                 data["descriptors1"] = (
235                     data["descriptors1"].transpose(-1, -2).contiguous()
236                 )
237                 runtime = measure(matcher, data, device=device, r=args.repeat)[
238                     "mean"
239                 ]
240                 results[pair_name][name].append(
241                     1000 / runtime if args.measure == "throughput" else runtime
242                 )
243             ax.plot(
244                 args.num_keypoints, results[pair_name][name], label=name, marker="o"
245             )
246         del matcher, data, image0, image1, feats0, feats1
247
248 for name, runtimes in results.items():
249     print_as_table(runtimes, name, args.num_keypoints)
250
251 axes[0].legend()
252 fig.tight_layout()

```

```
253     if args.save:
254         plt.savefig(args.save, dpi=fig.dpi)
255     plt.show()
```

```

0 LightGlue/demo.py
1 # If we are on colab: this clones the repo and installs the dependencies
2 from pathlib import Path
3
4 # if Path.cwd().name != "LightGlue":
5 #     !git clone --quiet https://github.com/cvg/LightGlue/
6 #     %cd LightGlue
7 #     !pip install --progress-bar off --quiet -e .
8
9 from lightglue import LightGlue, SuperPoint, DISK
10 from lightglue.utils import load_image, rbd
11 from lightglue import viz2d
12 import torch
13
14 torch.set_grad_enabled(False)
15 images = Path("assets")
16
17 device = torch.device("cuda" if torch.cuda.is_available() else "cpu") # 'mps', 'cpu'
18
19 extractor = SuperPoint(max_num_keypoints=2048).eval().to(device) # load the extractor
20 matcher = LightGlue(features="superpoint").eval().to(device)
21
22 image0 = load_image(images / "DSC_0411.JPG")
23 image1 = load_image(images / "DSC_0410.JPG")
24
25 feats0 = extractor.extract(image0.to(device))
26 feats1 = extractor.extract(image1.to(device))
27 matches01 = matcher({"image0": feats0, "image1": feats1})
28 feats0, feats1, matches01 = [
29     rbd(x) for x in [feats0, feats1, matches01]
30 ] # remove batch dimension
31
32 kpts0, kpts1, matches = feats0["keypoints"], feats1["keypoints"], matches01["matches"]
33 m_kpts0, m_kpts1 = kpts0[matches[:, 0]], kpts1[matches[:, 1]]
34
35 axes = viz2d.plot_images([image0, image1])
36 viz2d.plot_matches(m_kpts0, m_kpts1, color="lime", lw=0.2)
37 viz2d.add_text(0, f'Stop after {matches01["stop"]} layers', fs=20)
38
39 kpc0, kpc1 = viz2d.cm_prune(matches01["prune0"], viz2d.cm_prune(matches01["prune1"]))
40 viz2d.plot_images([image0, image1])
41 viz2d.plot_keypoints([kpts0, kpts1], colors=[kpc0, kpc1], ps=10)

```

```

0 LightGlue/lightglue/aliked.py
1 # BSD 3-Clause License
2
3 # Copyright (c) 2022, Zhao Xiaoming
4 # All rights reserved.
5
6 # Redistribution and use in source and binary forms, with or without
7 # modification, are permitted provided that the following conditions are met:
8
9 # 1. Redistributions of source code must retain the above copyright notice, this
10 #    list of conditions and the following disclaimer.
11
12 # 2. Redistributions in binary form must reproduce the above copyright notice,
13 #    this list of conditions and the following disclaimer in the documentation
14 #    and/or other materials provided with the distribution.
15
16 # 3. Neither the name of the copyright holder nor the names of its
17 #    contributors may be used to endorse or promote products derived from
18 #    this software without specific prior written permission.
19
20 # THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS"
21 # AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE
22 # IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE
23 # DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR CONTRIBUTORS BE LIABLE
24 # FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL
25 # DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR
26 # SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER
27 # CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY,
28 # OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
29 # OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
30
31 # Authors:
32 # Xiaoming Zhao, Xingming Wu, Weihai Chen, Peter C.Y. Chen, Qingsong Xu, and Zhengguo Li
33 # Code from https://github.com/Shiaoming/ALIKED
34
35 from typing import Callable, Optional
36
37 import torch
38 import torch.nn.functional as F
39 import torchvision
40 from kornia.color import grayscale_to_rgb
41 from torch import nn
42 from torch.nn.modules.utils import _pair
43 from torchvision.models import resnet
44
45 from .utils import Extractor
46
47
48 def get_patches(
49     tensor: torch.Tensor, required_corners: torch.Tensor, ps: int
50 ) -> torch.Tensor:
51     c, h, w = tensor.shape
52     corner = (required_corners - ps // 2 + 1).long()
53     corner[:, 0] = corner[:, 0].clamp(min=0, max=w - 1 - ps)
54     corner[:, 1] = corner[:, 1].clamp(min=0, max=h - 1 - ps)
55     offset = torch.arange(0, ps)
56
57     kw = {"indexing": "ij"} if torch.__version__ >= "1.10" else {}
58     x, y = torch.meshgrid(offset, offset, **kw)
59     patches = torch.stack((x, y)).permute(2, 1, 0).unsqueeze(2)
60     patches = patches.to(corner) + corner[None, None]
61     pts = patches.reshape(-1, 2)
62     sampled = tensor.permute(1, 2, 0)[tuple(pts.T)[:, :-1]]
63     sampled = sampled.reshape(ps, ps, -1, c)
64     assert sampled.shape[:3] == patches.shape[:3]
65     return sampled.permute(2, 3, 0, 1)
66
67
68 def simple_nms(scores: torch.Tensor, nms_radius: int):
69     """Fast Non-maximum suppression to remove nearby points"""
70
71     zeros = torch.zeros_like(scores)
72     max_mask = scores == torch.nn.functional.max_pool2d(
73         scores, kernel_size=nms_radius * 2 + 1, stride=1, padding=nms_radius
74     )
75
76     for _ in range(2):
77         supp_mask = (
78             torch.nn.functional.max_pool2d(
79                 max_mask.float(),
80                 kernel_size=nms_radius * 2 + 1,
81                 stride=1,
82                 padding=nms_radius,

```

```

83         )
84         > 0
85     )
86     supp_scores = torch.where(supp_mask, zeros, scores)
87     new_max_mask = supp_scores == torch.nn.functional.max_pool2d(
88         supp_scores, kernel_size=nms_radius * 2 + 1, stride=1, padding=nms_radius
89     )
90     max_mask = max_mask | (new_max_mask & (~supp_mask))
91     return torch.where(max_mask, scores, zeros)
92
93
94 class DKD(nn.Module):
95     def __init__(
96         self,
97         radius: int = 2,
98         top_k: int = 0,
99         scores_th: float = 0.2,
100         n_limit: int = 20000,
101     ):
102         """
103         Args:
104             radius: soft detection radius, kernel size is (2 * radius + 1)
105             top_k: top_k > 0: return top k keypoints
106             scores_th: top_k <= 0 threshold mode:
107                 scores_th > 0: return keypoints with scores>scores_th
108                 else: return keypoints with scores > scores.mean()
109             n_limit: max number of keypoint in threshold mode
110         """
111         super().__init__()
112         self.radius = radius
113         self.top_k = top_k
114         self.scores_th = scores_th
115         self.n_limit = n_limit
116         self.kernel_size = 2 * self.radius + 1
117         self.temperature = 0.1 # tuned temperature
118         self.unfold = nn.Unfold(kernel_size=self.kernel_size, padding=self.radius)
119         # local xy grid
120         x = torch.linspace(-self.radius, self.radius, self.kernel_size)
121         # (kernel_size*kernel_size) x 2 : (w,h)
122         kw = {"indexing": "ij"} if torch.__version__ >= "1.10" else {}
123         self.hw_grid = (
124             torch.stack(torch.meshgrid([x, x], **kw)).view(2, -1).t()[0, [1, 0]]
125         )
126
127     def forward(
128         self,
129         scores_map: torch.Tensor,
130         sub_pixel: bool = True,
131         image_size: Optional[torch.Tensor] = None,
132     ):
133         """
134         :param scores_map: Bx1xHxW
135         :param descriptor_map: BxCxHxW
136         :param sub_pixel: whether to use sub-pixel keypoint detection
137         :return: kpts: list[Nx2,...]; kptscores: list[N,...] normalised position: -1~1
138         """
139         b, c, h, w = scores_map.shape
140         scores_nograd = scores_map.detach()
141         nms_scores = simple_nms(scores_nograd, self.radius)
142
143         # remove border
144         nms_scores[:, :, : self.radius, :] = 0
145         nms_scores[:, :, :, : self.radius] = 0
146         if image_size is not None:
147             for i in range(scores_map.shape[0]):
148                 w, h = image_size[i].long()
149                 nms_scores[i, :, h.item() - self.radius :, :] = 0
150                 nms_scores[i, :, :, w.item() - self.radius :] = 0
151         else:
152             nms_scores[:, :, -self.radius :, :] = 0
153             nms_scores[:, :, :, -self.radius :] = 0
154
155         # detect keypoints without grad
156         if self.top_k > 0:
157             topk = torch.topk(nms_scores.view(b, -1), self.top_k)
158             indices_keypoints = [topk.indices[i] for i in range(b)] # B x top_k
159         else:
160             if self.scores_th > 0:
161                 masks = nms_scores > self.scores_th
162                 if masks.sum() == 0:
163                     th = scores_nograd.reshape(b, -1).mean(dim=1) # th = self.scores_th
164                     masks = nms_scores > th.reshape(b, 1, 1, 1)
165             else:
166                 th = scores_nograd.reshape(b, -1).mean(dim=1) # th = self.scores_th
167                 masks = nms_scores > th.reshape(b, 1, 1, 1)

```

```

168 masks = masks.reshape(b, -1)
169
170 indices_keypoints = [] # list, B x (any size)
171 scores_view = scores_nograd.reshape(b, -1)
172 for mask, scores in zip(masks, scores_view):
173     indices = mask.nonzero()[0]
174     if len(indices) > self.n_limit:
175         kpts_sc = scores[indices]
176         sort_idx = kpts_sc.sort(descending=True)[1]
177         sel_idx = sort_idx[: self.n_limit]
178         indices = indices[sel_idx]
179     indices_keypoints.append(indices)
180
181 wh = torch.tensor([w - 1, h - 1], device=scores_nograd.device)
182
183 keypoints = []
184 scoredispersitys = []
185 kptscores = []
186 if sub_pixel:
187     # detect soft keypoints with grad backpropagation
188     patches = self.unfold(scores_map) # B x (kernel**2) x (H*W)
189     self.hw_grid = self.hw_grid.to(scores_map) # to device
190     for b_idx in range(b):
191         patch = patches[b_idx].t() # (H*W) x (kernel**2)
192         indices_kpt = indices_keypoints[
193             b_idx
194         ] # one dimension vector, say its size is M
195         patch_scores = patch[indices_kpt] # M x (kernel**2)
196         keypoints_xy_nms = torch.stack(
197             [indices_kpt % w, torch.div(indices_kpt, w, rounding_mode="trunc")],
198             dim=1,
199         ) # Mx2
200
201         # max is detached to prevent undesired backprop loops in the graph
202         max_v = patch_scores.max(dim=1).values.detach()[0, None]
203         x_exp = (
204             (patch_scores - max_v) / self.temperature
205         ).exp() # M * (kernel**2), in [0, 1]
206
207         # \frac{\sum_{i,j} \times \exp(x/T)}{\sum \exp(x/T)}
208         xy_residual = (
209             x_exp @ self.hw_grid / x_exp.sum(dim=1)[0, None]
210         ) # Soft-argmax, Mx2
211
212         hw_grid_dist2 = (
213             torch.norm(
214                 (self.hw_grid[None, :, :] - xy_residual[:, None, :])
215                 / self.radius,
216                 dim=-1,
217             )
218             ** 2
219         )
220         scoredispersity = (x_exp * hw_grid_dist2).sum(dim=1) / x_exp.sum(dim=1)
221
222         # compute result keypoints
223         keypoints_xy = keypoints_xy_nms + xy_residual
224         keypoints_xy = keypoints_xy / wh * 2 - 1 # (w,h) -> (-1~1,-1~1)
225
226         kptscore = torch.nn.functional.grid_sample(
227             scores_map[b_idx].unsqueeze(0),
228             keypoints_xy.view(1, 1, -1, 2),
229             mode="bilinear",
230             align_corners=True,
231         )[
232             0, 0, 0, :
233         ] # CxN
234
235         keypoints.append(keypoints_xy)
236         scoredispersitys.append(scoredispersity)
237         kptscores.append(kptscore)
238 else:
239     for b_idx in range(b):
240         indices_kpt = indices_keypoints[
241             b_idx
242         ] # one dimension vector, say its size is M
243         # To avoid warning: UserWarning: __floordiv__ is deprecated
244         keypoints_xy_nms = torch.stack(
245             [indices_kpt % w, torch.div(indices_kpt, w, rounding_mode="trunc")],
246             dim=1,
247         ) # Mx2
248         keypoints_xy = keypoints_xy_nms / wh * 2 - 1 # (w,h) -> (-1~1,-1~1)
249         kptscore = torch.nn.functional.grid_sample(
250             scores_map[b_idx].unsqueeze(0),
251             keypoints_xy.view(1, 1, -1, 2),
252             mode="bilinear",

```



```

253         align_corners=True,
254     ) [
255         0, 0, 0, :
256     ] # CxN
257     keypoints.append(keypoints_xy)
258     scoredispersitys.append(kptscore) # for jit.script compatability
259     kptscores.append(kptscore)
260
261     return keypoints, scoredispersitys, kptscores
262
263
264 class InputPadder(object):
265     """Pads images such that dimensions are divisible by 8"""
266
267     def __init__(self, h: int, w: int, divis_by: int = 8):
268         self.ht = h
269         self.wd = w
270         pad_ht = (((self.ht // divis_by) + 1) * divis_by - self.ht) % divis_by
271         pad_wd = (((self.wd // divis_by) + 1) * divis_by - self.wd) % divis_by
272         self._pad = [
273             pad_wd // 2,
274             pad_wd - pad_wd // 2,
275             pad_ht // 2,
276             pad_ht - pad_ht // 2,
277         ]
278
279     def pad(self, x: torch.Tensor):
280         assert x.ndim == 4
281         return F.pad(x, self._pad, mode="replicate")
282
283     def unpad(self, x: torch.Tensor):
284         assert x.ndim == 4
285         ht = x.shape[-2]
286         wd = x.shape[-1]
287         c = [self._pad[2], ht - self._pad[3], self._pad[0], wd - self._pad[1]]
288         return x[..., c[0] : c[1], c[2] : c[3]]
289
290
291 class DeformableConv2d(nn.Module):
292     def __init__(
293         self,
294         in_channels,
295         out_channels,
296         kernel_size=3,
297         stride=1,
298         padding=1,
299         bias=False,
300         mask=False,
301     ):
302         super(DeformableConv2d, self).__init__()
303
304         self.padding = padding
305         self.mask = mask
306
307         self.channel_num = (
308             3 * kernel_size * kernel_size if mask else 2 * kernel_size * kernel_size
309         )
310         self.offset_conv = nn.Conv2d(
311             in_channels,
312             self.channel_num,
313             kernel_size=kernel_size,
314             stride=stride,
315             padding=self.padding,
316             bias=True,
317         )
318
319         self.regular_conv = nn.Conv2d(
320             in_channels=in_channels,
321             out_channels=out_channels,
322             kernel_size=kernel_size,
323             stride=stride,
324             padding=self.padding,
325             bias=bias,
326         )
327
328     def forward(self, x):
329         h, w = x.shape[2:]
330         max_offset = max(h, w) / 4.0
331
332         out = self.offset_conv(x)
333         if self.mask:
334             o1, o2, mask = torch.chunk(out, 3, dim=1)
335             offset = torch.cat((o1, o2), dim=1)
336             mask = torch.sigmoid(mask)
337         else:

```

```

338         offset = out
339         mask = None
340         offset = offset.clamp(-max_offset, max_offset)
341         x = torchvision.ops.deform_conv2d(
342             input=x,
343             offset=offset,
344             weight=self.regular_conv.weight,
345             bias=self.regular_conv.bias,
346             padding=self.padding,
347             mask=mask,
348         )
349         return x
350
351
352 def get_conv(
353     inplanes,
354     planes,
355     kernel_size=3,
356     stride=1,
357     padding=1,
358     bias=False,
359     conv_type="conv",
360     mask=False,
361 ):
362     if conv_type == "conv":
363         conv = nn.Conv2d(
364             inplanes,
365             planes,
366             kernel_size=kernel_size,
367             stride=stride,
368             padding=padding,
369             bias=bias,
370         )
371     elif conv_type == "dcn":
372         conv = DeformableConv2d(
373             inplanes,
374             planes,
375             kernel_size=kernel_size,
376             stride=stride,
377             padding=_pair(padding),
378             bias=bias,
379             mask=mask,
380         )
381     else:
382         raise TypeError
383     return conv
384
385
386 class ConvBlock(nn.Module):
387     def __init__(
388         self,
389         in_channels,
390         out_channels,
391         gate: Optional[Callable[..., nn.Module]] = None,
392         norm_layer: Optional[Callable[..., nn.Module]] = None,
393         conv_type: str = "conv",
394         mask: bool = False,
395     ):
396         super().__init__()
397         if gate is None:
398             self.gate = nn.ReLU(inplace=True)
399         else:
400             self.gate = gate
401         if norm_layer is None:
402             norm_layer = nn.BatchNorm2d
403         self.conv1 = get_conv(
404             in_channels, out_channels, kernel_size=3, conv_type=conv_type, mask=mask
405         )
406         self.bn1 = norm_layer(out_channels)
407         self.conv2 = get_conv(
408             out_channels, out_channels, kernel_size=3, conv_type=conv_type, mask=mask
409         )
410         self.bn2 = norm_layer(out_channels)
411
412     def forward(self, x):
413         x = self.gate(self.bn1(self.conv1(x))) # B x in_channels x H x W
414         x = self.gate(self.bn2(self.conv2(x))) # B x out_channels x H x W
415         return x
416
417
418 # modified based on torchvision\models\resnet.py#27->BasicBlock
419 class ResBlock(nn.Module):
420     expansion: int = 1
421
422     def __init__(

```

```

423         self,
424         inplanes: int,
425         planes: int,
426         stride: int = 1,
427         downsample: Optional[nn.Module] = None,
428         groups: int = 1,
429         base_width: int = 64,
430         dilation: int = 1,
431         gate: Optional[Callable[..., nn.Module]] = None,
432         norm_layer: Optional[Callable[..., nn.Module]] = None,
433         conv_type: str = "conv",
434         mask: bool = False,
435     ) -> None:
436         super(ResBlock, self).__init__()
437         if gate is None:
438             self.gate = nn.ReLU(inplace=True)
439         else:
440             self.gate = gate
441         if norm_layer is None:
442             norm_layer = nn.BatchNorm2d
443         if groups != 1 or base_width != 64:
444             raise ValueError("ResBlock only supports groups=1 and base_width=64")
445         if dilation > 1:
446             raise NotImplementedError("Dilation > 1 not supported in ResBlock")
447         # Both self.conv1 and self.downsample layers
448         # downsample the input when stride != 1
449         self.conv1 = get_conv(
450             inplanes, planes, kernel_size=3, conv_type=conv_type, mask=mask
451         )
452         self.bn1 = norm_layer(planes)
453         self.conv2 = get_conv(
454             planes, planes, kernel_size=3, conv_type=conv_type, mask=mask
455         )
456         self.bn2 = norm_layer(planes)
457         self.downsample = downsample
458         self.stride = stride
459
460     def forward(self, x: torch.Tensor) -> torch.Tensor:
461         identity = x
462
463         out = self.conv1(x)
464         out = self.bn1(out)
465         out = self.gate(out)
466
467         out = self.conv2(out)
468         out = self.bn2(out)
469
470         if self.downsample is not None:
471             identity = self.downsample(x)
472
473         out += identity
474         out = self.gate(out)
475
476         return out
477
478
479 class SDDH(nn.Module):
480     def __init__(
481         self,
482         dims: int,
483         kernel_size: int = 3,
484         n_pos: int = 8,
485         gate=nn.ReLU(),
486         conv2D=False,
487         mask=False,
488     ):
489         super(SDDH, self).__init__()
490         self.kernel_size = kernel_size
491         self.n_pos = n_pos
492         self.conv2D = conv2D
493         self.mask = mask
494
495         self.get_patches_func = get_patches
496
497         # estimate offsets
498         self.channel_num = 3 * n_pos if mask else 2 * n_pos
499         self.offset_conv = nn.Sequential(
500             nn.Conv2d(
501                 dims,
502                 self.channel_num,
503                 kernel_size=kernel_size,
504                 stride=1,
505                 padding=0,
506                 bias=True,
507             ),

```

```

508         gate,
509         nn.Conv2d(
510             self.channel_num,
511             self.channel_num,
512             kernel_size=1,
513             stride=1,
514             padding=0,
515             bias=True,
516         ),
517     )
518
519     # sampled feature conv
520     self.sf_conv = nn.Conv2d(
521         dims, dims, kernel_size=1, stride=1, padding=0, bias=False
522     )
523
524     # convM
525     if not conv2D:
526         # deformable desc weights
527         agg_weights = torch.nn.Parameter(torch.rand(n_pos, dims, dims))
528         self.register_parameter("agg_weights", agg_weights)
529     else:
530         self.convM = nn.Conv2d(
531             dims * n_pos, dims, kernel_size=1, stride=1, padding=0, bias=False
532         )
533
534     def forward(self, x, keypoints):
535         # x: [B,C,H,W]
536         # keypoints: list, [[N_kpts,2], ...] (w,h)
537         b, c, h, w = x.shape
538         wh = torch.tensor([w - 1, h - 1], device=x.device)
539         max_offset = max(h, w) / 4.0
540
541         offsets = []
542         descriptors = []
543         # get offsets for each keypoint
544         for ib in range(b):
545             xi, kptsi = x[ib], keypoints[ib]
546             kptsi_wh = (kptsi / 2 + 0.5) * wh
547             N_kpts = len(kptsi)
548
549             if self.kernel_size > 1:
550                 patch = self.get_patches_func(
551                     xi, kptsi_wh.long(), self.kernel_size
552                 ) # [N_kpts, C, K, K]
553             else:
554                 kptsi_wh_long = kptsi_wh.long()
555                 patch = (
556                     xi[:, kptsi_wh_long[:, 1], kptsi_wh_long[:, 0]]
557                     .permute(1, 0)
558                     .reshape(N_kpts, c, 1, 1)
559                 )
560
561             offset = self.offset_conv(patch).clamp(
562                 -max_offset, max_offset
563             ) # [N_kpts, 2*n_pos, 1, 1]
564             if self.mask:
565                 offset = (
566                     offset[:, :, 0, 0].view(N_kpts, 3, self.n_pos).permute(0, 2, 1)
567                 ) # [N_kpts, n_pos, 3]
568                 offset = offset[:, :, :-1] # [N_kpts, n_pos, 2]
569                 mask_weight = torch.sigmoid(offset[:, :, -1]) # [N_kpts, n_pos]
570             else:
571                 offset = (
572                     offset[:, :, 0, 0].view(N_kpts, 2, self.n_pos).permute(0, 2, 1)
573                 ) # [N_kpts, n_pos, 2]
574             offsets.append(offset) # for visualization
575
576             # get sample positions
577             pos = kptsi_wh.unsqueeze(1) + offset # [N_kpts, n_pos, 2]
578             pos = 2.0 * pos / wh[None] - 1
579             pos = pos.reshape(1, N_kpts * self.n_pos, 1, 2)
580
581             # sample features
582             features = F.grid_sample(
583                 xi.unsqueeze(0), pos, mode="bilinear", align_corners=True
584             ) # [1,C,(N_kpts*n_pos),1]
585             features = features.reshape(c, N_kpts, self.n_pos, 1).permute(
586                 1, 0, 2, 3
587             ) # [N_kpts, C, n_pos, 1]
588             if self.mask:
589                 features = torch.einsum("ncpo,np->ncpo", features, mask_weight)
590
591             features = torch.selu_(self.sf_conv(features)).squeeze(
592                 -1

```

```

593         ) # [N_kpts, C, n_pos]
594         # convM
595         if not self.conv2D:
596             descs = torch.einsum(
597                 "ncp,pcd->nd", features, self.agg_weights
598             ) # [N_kpts, C]
599         else:
600             features = features.reshape(N_kpts, -1)[
601                 :, :, None, None
602             ] # [N_kpts, C*n_pos, 1, 1]
603             descs = self.convM(features).squeeze() # [N_kpts, C]
604
605         # normalize
606         descs = F.normalize(descs, p=2.0, dim=1)
607         descriptors.append(descs)
608
609     return descriptors, offsets
610
611
612 class ALIKED(Extractor):
613     default_conf = {
614         "model_name": "aliked-nl6",
615         "max_num_keypoints": -1,
616         "detection_threshold": 0.2,
617         "nms_radius": 2,
618     }
619
620     checkpoint_url = "https://github.com/Shiaoming/ALIKED/raw/main/models/{}.pth"
621
622     n_limit_max = 20000
623
624     # c1, c2, c3, c4, dim, K, M
625     cfgs = {
626         "aliked-t16": [8, 16, 32, 64, 64, 3, 16],
627         "aliked-nl6": [16, 32, 64, 128, 128, 3, 16],
628         "aliked-nl6rot": [16, 32, 64, 128, 128, 3, 16],
629         "aliked-n32": [16, 32, 64, 128, 128, 3, 32],
630     }
631
632     preprocess_conf = {
633         "resize": 1024,
634     }
635
636     required_data_keys = ["image"]
637
638     def __init__(self, **conf):
639         super().__init__(**conf) # Update with default configuration.
640         conf = self.conf
641         c1, c2, c3, c4, dim, K, M = self.cfgs[conf.model_name]
642         conv_types = ["conv", "conv", "dcn", "dcn"]
643         conv2D = False
644         mask = False
645
646         # build model
647         self.pool2 = nn.AvgPool2d(kernel_size=2, stride=2)
648         self.pool4 = nn.AvgPool2d(kernel_size=4, stride=4)
649         self.norm = nn.BatchNorm2d
650         self.gate = nn.SELU(inplace=True)
651         self.block1 = ConvBlock(3, c1, self.gate, self.norm, conv_type=conv_types[0])
652         self.block2 = self.get_resblock(c1, c2, conv_types[1], mask)
653         self.block3 = self.get_resblock(c2, c3, conv_types[2], mask)
654         self.block4 = self.get_resblock(c3, c4, conv_types[3], mask)
655
656         self.conv1 = resnet.conv1x1(c1, dim // 4)
657         self.conv2 = resnet.conv1x1(c2, dim // 4)
658         self.conv3 = resnet.conv1x1(c3, dim // 4)
659         self.conv4 = resnet.conv1x1(dim, dim // 4)
660         self.upsample2 = nn.Upsample(
661             scale_factor=2, mode="bilinear", align_corners=True
662         )
663         self.upsample4 = nn.Upsample(
664             scale_factor=4, mode="bilinear", align_corners=True
665         )
666         self.upsample8 = nn.Upsample(
667             scale_factor=8, mode="bilinear", align_corners=True
668         )
669         self.upsample32 = nn.Upsample(
670             scale_factor=32, mode="bilinear", align_corners=True
671         )
672         self.score_head = nn.Sequential(
673             resnet.conv1x1(dim, 8),
674             self.gate,
675             resnet.conv3x3(8, 4),
676             self.gate,
677             resnet.conv3x3(4, 4),
678             self.gate,

```

```

678         resnet.conv3x3(4, 1),
679     )
680     self.desc_head = SDDH(dim, K, M, gate=self.gate, conv2D=conv2D, mask=mask)
681     self.dkd = DKD(
682         radius=conf.nms_radius,
683         top_k=-1 if conf.detection_threshold > 0 else conf.max_num_keypoints,
684         scores_th=conf.detection_threshold,
685         n_limit=conf.max_num_keypoints
686         if conf.max_num_keypoints > 0
687         else self.n_limit_max,
688     )
689
690     state_dict = torch.hub.load_state_dict_from_url(
691         self.checkpoint_url.format(conf.model_name), map_location="cpu"
692     )
693     self.load_state_dict(state_dict, strict=True)
694
695     def get_resblock(self, c_in, c_out, conv_type, mask):
696         return ResBlock(
697             c_in,
698             c_out,
699             1,
700             nn.Conv2d(c_in, c_out, 1),
701             gate=self.gate,
702             norm_layer=self.norm,
703             conv_type=conv_type,
704             mask=mask,
705         )
706
707     def extract_dense_map(self, image):
708         # Pads images such that dimensions are divisible by
709         div_by = 2**5
710         padder = InputPadder(image.shape[-2], image.shape[-1], div_by)
711         image = padder.pad(image)
712
713         # ===== feature encoder
714         x1 = self.block1(image) # B x c1 x H x W
715         x2 = self.pool2(x1)
716         x2 = self.block2(x2) # B x c2 x H/2 x W/2
717         x3 = self.pool4(x2)
718         x3 = self.block3(x3) # B x c3 x H/8 x W/8
719         x4 = self.pool4(x3)
720         x4 = self.block4(x4) # B x dim x H/32 x W/32
721         # ===== feature aggregation
722         x1 = self.gate(self.conv1(x1)) # B x dim//4 x H x W
723         x2 = self.gate(self.conv2(x2)) # B x dim//4 x H//2 x W//2
724         x3 = self.gate(self.conv3(x3)) # B x dim//4 x H//8 x W//8
725         x4 = self.gate(self.conv4(x4)) # B x dim//4 x H//32 x W//32
726         x2_up = self.upsample2(x2) # B x dim//4 x H x W
727         x3_up = self.upsample8(x3) # B x dim//4 x H x W
728         x4_up = self.upsample32(x4) # B x dim//4 x H x W
729         x1234 = torch.cat([x1, x2_up, x3_up, x4_up], dim=1)
730         # ===== score head
731         score_map = torch.sigmoid(self.score_head(x1234))
732         feature_map = torch.nn.functional.normalize(x1234, p=2, dim=1)
733
734         # Unpads images
735         feature_map = padder.unpad(feature_map)
736         score_map = padder.unpad(score_map)
737
738         return feature_map, score_map
739
740     def forward(self, data: dict) -> dict:
741         image = data["image"]
742         if image.shape[1] == 1:
743             image = grayscale_to_rgb(image)
744         feature_map, score_map = self.extract_dense_map(image)
745         keypoints, kptscores, scoredispersitys = self.dkd(
746             score_map, image_size=data.get("image_size")
747         )
748         descriptors, offsets = self.desc_head(feature_map, keypoints)
749
750         _, _, h, w = image.shape
751         wh = torch.tensor([w - 1, h - 1], device=image.device)
752         # no padding required
753         # we can set detection_threshold=-1 and conf.max_num_keypoints > 0
754         return {
755             "keypoints": wh * (torch.stack(keypoints) + 1) / 2.0, # B x N x 2
756             "descriptors": torch.stack(descriptors), # B x N x D
757             "keypoint_scores": torch.stack(kptscores), # B x N
758         }

```

```

0 LightGlue/lightglue/disk.py
1 import kornia
2 import torch
3
4 from .utils import Extractor
5
6
7 class DISK(Extractor):
8     default_conf = {
9         "weights": "depth",
10         "max_num_keypoints": None,
11         "desc_dim": 128,
12         "nms_window_size": 5,
13         "detection_threshold": 0.0,
14         "pad_if_not_divisible": True,
15     }
16
17     preprocess_conf = {
18         "resize": 1024,
19         "grayscale": False,
20     }
21
22     required_data_keys = ["image"]
23
24     def __init__(self, **conf) -> None:
25         super().__init__(**conf) # Update with default configuration.
26         self.model = kornia.feature.DISK.from_pretrained(self.conf.weights)
27
28     def forward(self, data: dict) -> dict:
29         """Compute keypoints, scores, descriptors for image"""
30         for key in self.required_data_keys:
31             assert key in data, f"Missing key {key} in data"
32         image = data["image"]
33         if image.shape[1] == 1:
34             image = kornia.color.grayscale_to_rgb(image)
35         features = self.model(
36             image,
37             n=self.conf.max_num_keypoints,
38             window_size=self.conf.nms_window_size,
39             score_threshold=self.conf.detection_threshold,
40             pad_if_not_divisible=self.conf.pad_if_not_divisible,
41         )
42         keypoints = [f.keypoints for f in features]
43         scores = [f.detection_scores for f in features]
44         descriptors = [f.descriptors for f in features]
45         del features
46
47         keypoints = torch.stack(keypoints, 0)
48         scores = torch.stack(scores, 0)
49         descriptors = torch.stack(descriptors, 0)
50
51         return {
52             "keypoints": keypoints.to(image).contiguous(),
53             "keypoint_scores": scores.to(image).contiguous(),
54             "descriptors": descriptors.to(image).contiguous(),
55         }

```

```

0 LightGlue/lightglue/dog_hardnet.py
1 import torch
2 from kornia.color import rgb_to_grayscale
3 from kornia.feature import HardNet, LAFDescriptor, laf_from_center_scale_ori
4
5 from .sift import SIFT
6
7
8 class DoGHardNet(SIFT):
9     required_data_keys = ["image"]
10
11     def __init__(self, **conf):
12         super().__init__(**conf)
13         self.laf_desc = LAFDescriptor(HardNet(True)).eval()
14
15     def forward(self, data: dict) -> dict:
16         image = data["image"]
17         if image.shape[1] == 3:
18             image = rgb_to_grayscale(image)
19         device = image.device
20         self.laf_desc = self.laf_desc.to(device)
21         self.laf_desc.descriptor = self.laf_desc.descriptor.eval()
22         pred = []
23         if "image_size" in data.keys():
24             im_size = data.get("image_size").long()
25         else:
26             im_size = None
27         for k in range(len(image)):
28             img = image[k]
29             if im_size is not None:
30                 w, h = data["image_size"][k]
31                 img = img[:, :h.to(torch.int32), :w.to(torch.int32)]
32             p = self.extract_single_image(img)
33             lafs = laf_from_center_scale_ori(
34                 p["keypoints"].reshape(1, -1, 2),
35                 6.0 * p["scales"].reshape(1, -1, 1, 1),
36                 torch.rad2deg(p["oris"]).reshape(1, -1, 1),
37             ).to(device)
38             p["descriptors"] = self.laf_desc(img[None], lafs).reshape(-1, 128)
39             pred.append(p)
40         pred = {k: torch.stack([p[k] for p in pred], 0).to(device) for k in pred[0]}
41         return pred

```



```

0 LightGlue/lightglue/lightglue.py
1 import warnings
2 from pathlib import Path
3 from types import SimpleNamespace
4 from typing import Callable, List, Optional, Tuple
5
6 import numpy as np
7 import torch
8 import torch.nn.functional as F
9 from torch import nn
10
11 try:
12     from flash_attn.modules.mha import FlashCrossAttention
13 except ModuleNotFoundError:
14     FlashCrossAttention = None
15
16 if FlashCrossAttention or hasattr(F, "scaled_dot_product_attention"):
17     FLASH_AVAILABLE = True
18 else:
19     FLASH_AVAILABLE = False
20
21 torch.backends.cudnn.deterministic = True
22
23
24 @torch.cuda.amp.custom_fwd(cast_inputs=torch.float32)
25 def normalize_keypoints(
26     kpts: torch.Tensor, size: Optional[torch.Tensor] = None
27 ) -> torch.Tensor:
28     if size is None:
29         size = 1 + kpts.max(-2).values - kpts.min(-2).values
30     elif not isinstance(size, torch.Tensor):
31         size = torch.tensor(size, device=kpts.device, dtype=kpts.dtype)
32     size = size.to(kpts)
33     shift = size / 2
34     scale = size.max(-1).values / 2
35     kpts = (kpts - shift[..., None, :]) / scale[..., None, None]
36     return kpts
37
38
39 def pad_to_length(x: torch.Tensor, length: int) -> Tuple[torch.Tensor]:
40     if length <= x.shape[-2]:
41         return x, torch.ones_like(x[..., :1], dtype=torch.bool)
42     pad = torch.ones(
43         *x.shape[:-2], length - x.shape[-2], x.shape[-1], device=x.device, dtype=x.dtype
44     )
45     y = torch.cat([x, pad], dim=-2)
46     mask = torch.zeros(*y.shape[:-1], 1, dtype=torch.bool, device=x.device)
47     mask[..., : x.shape[-2], :] = True
48     return y, mask
49
50
51 def rotate_half(x: torch.Tensor) -> torch.Tensor:
52     x = x.unflatten(-1, (-1, 2))
53     x1, x2 = x.unbind(dim=-1)
54     return torch.stack((-x2, x1), dim=-1).flatten(start_dim=-2)
55
56
57 def apply_cached_rotary_emb(freqs: torch.Tensor, t: torch.Tensor) -> torch.Tensor:
58     return (t * freqs[0]) + (rotate_half(t) * freqs[1])
59
60
61 class LearnableFourierPositionalEncoding(nn.Module):
62     def __init__(self, M: int, dim: int, F_dim: int = None, gamma: float = 1.0) -> None:
63         super().__init__()
64         F_dim = F_dim if F_dim is not None else dim
65         self.gamma = gamma
66         self.Wr = nn.Linear(M, F_dim // 2, bias=False)
67         nn.init.normal_(self.Wr.weight.data, mean=0, std=self.gamma**-2)
68
69     def forward(self, x: torch.Tensor) -> torch.Tensor:
70         """encode position vector"""
71         projected = self.Wr(x)
72         cosines, sines = torch.cos(projected), torch.sin(projected)
73         emb = torch.stack([cosines, sines], 0).unsqueeze(-3)
74         return emb.repeat_interleave(2, dim=-1)
75
76
77 class TokenConfidence(nn.Module):
78     def __init__(self, dim: int) -> None:
79         super().__init__()
80         self.token = nn.Sequential(nn.Linear(dim, 1), nn.Sigmoid())
81
82     def forward(self, desc0: torch.Tensor, desc1: torch.Tensor):

```

```

83         """get confidence tokens"""
84     return (
85         self.token(desc0.detach()).squeeze(-1),
86         self.token(desc1.detach()).squeeze(-1),
87     )
88
89
90 class Attention(nn.Module):
91     def __init__(self, allow_flash: bool) -> None:
92         super().__init__()
93         if allow_flash and not FLASH_AVAILABLE:
94             warnings.warn(
95                 "FlashAttention is not available. For optimal speed, "
96                 "consider installing torch >= 2.0 or flash-attn.",
97                 stacklevel=2,
98             )
99         self.enable_flash = allow_flash and FLASH_AVAILABLE
100         self.has_sdp = hasattr(F, "scaled_dot_product_attention")
101         if allow_flash and FlashCrossAttention:
102             self.flash_ = FlashCrossAttention()
103         if self.has_sdp:
104             torch.backends.cuda.enable_flash_sdp(allow_flash)
105
106     def forward(self, q, k, v, mask: Optional[torch.Tensor] = None) -> torch.Tensor:
107         if q.shape[-2] == 0 or k.shape[-2] == 0:
108             return q.new_zeros((*q.shape[:-1], v.shape[-1]))
109         if self.enable_flash and q.device.type == "cuda":
110             # use torch 2.0 scaled_dot_product_attention with flash
111             if self.has_sdp:
112                 args = [x.half().contiguous() for x in [q, k, v]]
113                 v = F.scaled_dot_product_attention(*args, attn_mask=mask).to(q.dtype)
114                 return v if mask is None else v.nan_to_num()
115             else:
116                 assert mask is None
117                 q, k, v = [x.transpose(-2, -3).contiguous() for x in [q, k, v]]
118                 m = self.flash_(q.half(), torch.stack([k, v], 2).half())
119                 return m.transpose(-2, -3).to(q.dtype).clone()
120         elif self.has_sdp:
121             args = [x.contiguous() for x in [q, k, v]]
122             v = F.scaled_dot_product_attention(*args, attn_mask=mask)
123             return v if mask is None else v.nan_to_num()
124         else:
125             s = q.shape[-1] ** -0.5
126             sim = torch.einsum("...id,...jd->...ij", q, k) * s
127             if mask is not None:
128                 sim.masked_fill(~mask, -float("inf"))
129             attn = F.softmax(sim, -1)
130             return torch.einsum("...ij,...jd->...id", attn, v)
131
132
133 class SelfBlock(nn.Module):
134     def __init__(
135         self, embed_dim: int, num_heads: int, flash: bool = False, bias: bool = True
136     ) -> None:
137         super().__init__()
138         self.embed_dim = embed_dim
139         self.num_heads = num_heads
140         assert self.embed_dim % num_heads == 0
141         self.head_dim = self.embed_dim // num_heads
142         self.Wqkv = nn.Linear(embed_dim, 3 * embed_dim, bias=bias)
143         self.inner_attn = Attention(flash)
144         self.out_proj = nn.Linear(embed_dim, embed_dim, bias=bias)
145         self.ffn = nn.Sequential(
146             nn.Linear(2 * embed_dim, 2 * embed_dim),
147             nn.LayerNorm(2 * embed_dim, elementwise_affine=True),
148             nn.GELU(),
149             nn.Linear(2 * embed_dim, embed_dim),
150         )
151
152     def forward(
153         self,
154         x: torch.Tensor,
155         encoding: torch.Tensor,
156         mask: Optional[torch.Tensor] = None,
157     ) -> torch.Tensor:
158         qkv = self.Wqkv(x)
159         qkv = qkv.unflatten(-1, (self.num_heads, -1, 3)).transpose(1, 2)
160         q, k, v = qkv[..., 0], qkv[..., 1], qkv[..., 2]
161         q = apply_cached_rotary_emb(encoding, q)
162         k = apply_cached_rotary_emb(encoding, k)
163         context = self.inner_attn(q, k, v, mask=mask)
164         message = self.out_proj(context.transpose(1, 2).flatten(start_dim=-2))
165         return x + self.ffn(torch.cat([x, message], -1))
166
167

```

```

168 class CrossBlock(nn.Module):
169     def __init__(
170         self, embed_dim: int, num_heads: int, flash: bool = False, bias: bool = True
171     ) -> None:
172         super().__init__()
173         self.heads = num_heads
174         dim_head = embed_dim // num_heads
175         self.scale = dim_head**-0.5
176         inner_dim = dim_head * num_heads
177         self.to_qk = nn.Linear(embed_dim, inner_dim, bias=bias)
178         self.to_v = nn.Linear(embed_dim, inner_dim, bias=bias)
179         self.to_out = nn.Linear(inner_dim, embed_dim, bias=bias)
180         self.ffn = nn.Sequential(
181             nn.Linear(2 * embed_dim, 2 * embed_dim),
182             nn.LayerNorm(2 * embed_dim, elementwise_affine=True),
183             nn.GELU(),
184             nn.Linear(2 * embed_dim, embed_dim),
185         )
186         if flash and FLASH_AVAILABLE:
187             self.flash = Attention(True)
188         else:
189             self.flash = None
190
191     def map(self, func: Callable, x0: torch.Tensor, x1: torch.Tensor):
192         return func(x0), func(x1)
193
194     def forward(
195         self, x0: torch.Tensor, x1: torch.Tensor, mask: Optional[torch.Tensor] = None
196     ) -> List[torch.Tensor]:
197         qk0, qk1 = self.map(self.to_qk, x0, x1)
198         v0, v1 = self.map(self.to_v, x0, x1)
199         qk0, qk1, v0, v1 = map(
200             lambda t: t.unflatten(-1, (self.heads, -1)).transpose(1, 2),
201             (qk0, qk1, v0, v1),
202         )
203         if self.flash is not None and qk0.device.type == "cuda":
204             m0 = self.flash(qk0, qk1, v1, mask)
205             m1 = self.flash(
206                 qk1, qk0, v0, mask.transpose(-1, -2) if mask is not None else None
207             )
208         else:
209             qk0, qk1 = qk0 * self.scale**0.5, qk1 * self.scale**0.5
210             sim = torch.einsum("bhij, bhjd -> bhij", qk0, qk1)
211             if mask is not None:
212                 sim = sim.masked_fill(~mask, -float("inf"))
213             attn01 = F.softmax(sim, dim=-1)
214             attn10 = F.softmax(sim.transpose(-2, -1).contiguous(), dim=-1)
215             m0 = torch.einsum("bhij, bhjd -> bhid", attn01, v1)
216             m1 = torch.einsum("bhji, bhjd -> bhid", attn10.transpose(-2, -1), v0)
217             if mask is not None:
218                 m0, m1 = m0.nan_to_num(), m1.nan_to_num()
219             m0, m1 = self.map(lambda t: t.transpose(1, 2).flatten(start_dim=-2), m0, m1)
220             m0, m1 = self.map(self.to_out, m0, m1)
221             x0 = x0 + self.ffn(torch.cat([x0, m0], -1))
222             x1 = x1 + self.ffn(torch.cat([x1, m1], -1))
223             return x0, x1
224
225
226 class TransformerLayer(nn.Module):
227     def __init__(self, *args, **kwargs):
228         super().__init__()
229         self.self_attn = SelfBlock(*args, **kwargs)
230         self.cross_attn = CrossBlock(*args, **kwargs)
231
232     def forward(
233         self,
234         desc0,
235         desc1,
236         encoding0,
237         encoding1,
238         mask0: Optional[torch.Tensor] = None,
239         mask1: Optional[torch.Tensor] = None,
240     ):
241         if mask0 is not None and mask1 is not None:
242             return self.masked_forward(desc0, desc1, encoding0, encoding1, mask0, mask1)
243         else:
244             desc0 = self.self_attn(desc0, encoding0)
245             desc1 = self.self_attn(desc1, encoding1)
246             return self.cross_attn(desc0, desc1)
247
248     # This part is compiled and allows padding inputs
249     def masked_forward(self, desc0, desc1, encoding0, encoding1, mask0, mask1):
250         mask = mask0 & mask1.transpose(-1, -2)
251         mask0 = mask0 & mask0.transpose(-1, -2)
252         mask1 = mask1 & mask1.transpose(-1, -2)

```

```

253         desc0 = self.self_attn(desc0, encoding0, mask0)
254         desc1 = self.self_attn(desc1, encoding1, mask1)
255         return self.cross_attn(desc0, desc1, mask)
256
257
258 def sigmoid_log_double_softmax(
259     sim: torch.Tensor, z0: torch.Tensor, z1: torch.Tensor
260 ) -> torch.Tensor:
261     """create the log assignment matrix from logits and similarity"""
262     b, m, n = sim.shape
263     certainties = F.logsigmoid(z0) + F.logsigmoid(z1).transpose(1, 2)
264     scores0 = F.log_softmax(sim, 2)
265     scores1 = F.log_softmax(sim.transpose(-1, -2).contiguous(), 2).transpose(-1, -2)
266     scores = sim.new_full((b, m + 1, n + 1), 0)
267     scores[:, :m, :n] = scores0 + scores1 + certainties
268     scores[:, :-1, -1] = F.logsigmoid(-z0.squeeze(-1))
269     scores[:, -1, :-1] = F.logsigmoid(-z1.squeeze(-1))
270     return scores
271
272
273 class MatchAssignment(nn.Module):
274     def __init__(self, dim: int) -> None:
275         super().__init__()
276         self.dim = dim
277         self.matchability = nn.Linear(dim, 1, bias=True)
278         self.final_proj = nn.Linear(dim, dim, bias=True)
279
280     def forward(self, desc0: torch.Tensor, desc1: torch.Tensor):
281         """build assignment matrix from descriptors"""
282         mdesc0, mdesc1 = self.final_proj(desc0), self.final_proj(desc1)
283         _, _, d = mdesc0.shape
284         mdesc0, mdesc1 = mdesc0 / d**0.25, mdesc1 / d**0.25
285         sim = torch.einsum("bmd,bnd->bmn", mdesc0, mdesc1)
286         z0 = self.matchability(desc0)
287         z1 = self.matchability(desc1)
288         scores = sigmoid_log_double_softmax(sim, z0, z1)
289         return scores, sim
290
291     def get_matchability(self, desc: torch.Tensor):
292         return torch.sigmoid(self.matchability(desc)).squeeze(-1)
293
294
295 def filter_matches(scores: torch.Tensor, th: float):
296     """obtain matches from a log assignment matrix [Bx M+1 x N+1]"""
297     max0, max1 = scores[:, :-1, :-1].max(2), scores[:, :-1, :-1].max(1)
298     m0, m1 = max0.indices, max1.indices
299     indices0 = torch.arange(m0.shape[1], device=m0.device)[None]
300     indices1 = torch.arange(m1.shape[1], device=m1.device)[None]
301     mutual0 = indices0 == m1.gather(1, m0)
302     mutual1 = indices1 == m0.gather(1, m1)
303     max0_exp = max0.values.exp()
304     zero = max0_exp.new_tensor(0)
305     mscores0 = torch.where(mutual0, max0_exp, zero)
306     mscores1 = torch.where(mutual1, mscores0.gather(1, m1), zero)
307     valid0 = mutual0 & (mscores0 > th)
308     valid1 = mutual1 & valid0.gather(1, m1)
309     m0 = torch.where(valid0, m0, -1)
310     m1 = torch.where(valid1, m1, -1)
311     return m0, m1, mscores0, mscores1
312
313
314 class LightGlue(nn.Module):
315     default_conf = {
316         "name": "lightglue", # just for interfacing
317         "input_dim": 256, # input descriptor dimension (autoselected from weights)
318         "descriptor_dim": 256,
319         "add_scale_ori": False,
320         "n_layers": 9,
321         "num_heads": 4,
322         "flash": True, # enable FlashAttention if available.
323         "mp": False, # enable mixed precision
324         "depth_confidence": 0.95, # early stopping, disable with -1
325         "width_confidence": 0.99, # point pruning, disable with -1
326         "filter_threshold": 0.1, # match threshold
327         "weights": None,
328     }
329
330     # Point pruning involves an overhead (gather).
331     # Therefore, we only activate it if there are enough keypoints.
332     pruning_keypoint_thresholds = {
333         "cpu": -1,
334         "mps": -1,
335         "cuda": 1024,
336         "flash": 1536,
337     }

```

```

338
339 required_data_keys = ["image0", "image1"]
340
341 version = "v0.1_arxiv"
342 url = "https://github.com/cvg/LightGlue/releases/download/{}/{}_lightglue.pth"
343
344 features = {
345     "superpoint": {
346         "weights": "superpoint_lightglue",
347         "input_dim": 256,
348     },
349     "disk": {
350         "weights": "disk_lightglue",
351         "input_dim": 128,
352     },
353     "aliked": {
354         "weights": "aliked_lightglue",
355         "input_dim": 128,
356     },
357     "sift": {
358         "weights": "sift_lightglue",
359         "input_dim": 128,
360         "add_scale_ori": True,
361     },
362     "doghardnet": {
363         "weights": "doghardnet_lightglue",
364         "input_dim": 128,
365         "add_scale_ori": True,
366     },
367 }
368
369 def __init__(self, features="superpoint", **conf) -> None:
370     super().__init__()
371     self.conf = conf = SimpleNamespace(**{**self.default_conf, **conf})
372     if features is not None:
373         if features not in self.features:
374             raise ValueError(
375                 f"Unsupported features: {features} not in "
376                 f"{', '.join(self.features)}"
377             )
378         for k, v in self.features[features].items():
379             setattr(conf, k, v)
380
381     if conf.input_dim != conf.descriptor_dim:
382         self.input_proj = nn.Linear(conf.input_dim, conf.descriptor_dim, bias=True)
383     else:
384         self.input_proj = nn.Identity()
385
386     head_dim = conf.descriptor_dim // conf.num_heads
387     self.posenc = LearnableFourierPositionalEncoding(
388         2 + 2 * self.conf.add_scale_ori, head_dim, head_dim
389     )
390
391     h, n, d = conf.num_heads, conf.n_layers, conf.descriptor_dim
392
393     self.transformers = nn.ModuleList(
394         [TransformerLayer(d, h, conf.flash) for _ in range(n)]
395     )
396
397     self.log_assignment = nn.ModuleList([MatchAssignment(d) for _ in range(n)])
398     self.token_confidence = nn.ModuleList(
399         [TokenConfidence(d) for _ in range(n - 1)]
400     )
401     self.register_buffer(
402         "confidence_thresholds",
403         torch.Tensor(
404             [self.confidence_threshold(i) for i in range(self.conf.n_layers)]
405         ),
406     )
407
408     state_dict = None
409     if features is not None:
410         fname = f"{conf.weights}_{self.version.replace('.', '-')}.pth"
411         state_dict = torch.hub.load_state_dict_from_url(
412             self.url.format(self.version, features), file_name=fname
413         )
414         self.load_state_dict(state_dict, strict=False)
415     elif conf.weights is not None:
416         path = Path(__file__).parent
417         path = path / "weights/{}".format(self.conf.weights)
418         state_dict = torch.load(str(path), map_location="cpu")
419
420     if state_dict:
421         # rename old state dict entries
422         for i in range(self.conf.n_layers):

```

```

423         pattern = f"self_attn.{i}", f"transformers.{i}.self_attn"
424         state_dict = {k.replace(*pattern): v for k, v in state_dict.items()}
425         pattern = f"cross_attn.{i}", f"transformers.{i}.cross_attn"
426         state_dict = {k.replace(*pattern): v for k, v in state_dict.items()}
427         self.load_state_dict(state_dict, strict=False)
428
429     # static lengths LightGlue is compiled for (only used with torch.compile)
430     self.static_lengths = None
431
432     def compile(
433         self, mode="reduce-overhead", static_lengths=[256, 512, 768, 1024, 1280, 1536]
434     ):
435         if self.conf.width_confidence != -1:
436             warnings.warn(
437                 "Point pruning is partially disabled for compiled forward.",
438                 stacklevel=2,
439             )
440
441         torch._inductor.cudagraph_mark_step_begin()
442         for i in range(self.conf.n_layers):
443             self.transformers[i].masked_forward = torch.compile(
444                 self.transformers[i].masked_forward, mode=mode, fullgraph=True
445             )
446
447         self.static_lengths = static_lengths
448
449     def forward(self, data: dict) -> dict:
450         """
451         Match keypoints and descriptors between two images
452
453         Input (dict):
454             image0: dict
455                 keypoints: [B x M x 2]
456                 descriptors: [B x M x D]
457                 image: [B x C x H x W] or image_size: [B x 2]
458             image1: dict
459                 keypoints: [B x N x 2]
460                 descriptors: [B x N x D]
461                 image: [B x C x H x W] or image_size: [B x 2]
462         Output (dict):
463             matches0: [B x M]
464             matching_scores0: [B x M]
465             matches1: [B x N]
466             matching_scores1: [B x N]
467             matches: List[[Si x 2]]
468             scores: List[[Si]]
469             stop: int
470             prune0: [B x M]
471             pruned1: [B x N]
472         """
473         with torch.autocast(enabled=self.conf.mp, device_type="cuda"):
474             return self._forward(data)
475
476     def _forward(self, data: dict) -> dict:
477         for key in self.required_data_keys:
478             assert key in data, f"Missing key {key} in data"
479         data0, data1 = data["image0"], data["image1"]
480         kpts0, kpts1 = data0["keypoints"], data1["keypoints"]
481         b, m, _ = kpts0.shape
482         b, n, _ = kpts1.shape
483         device = kpts0.device
484         size0, size1 = data0.get("image_size"), data1.get("image_size")
485         kpts0 = normalize_keypoints(kpts0, size0).clone()
486         kpts1 = normalize_keypoints(kpts1, size1).clone()
487
488         if self.conf.add_scale_ori:
489             kpts0 = torch.cat(
490                 [kpts0] + [data0[k].unsqueeze(-1) for k in ("scales", "oris")], -1
491             )
492             kpts1 = torch.cat(
493                 [kpts1] + [data1[k].unsqueeze(-1) for k in ("scales", "oris")], -1
494             )
495         desc0 = data0["descriptors"].detach().contiguous()
496         desc1 = data1["descriptors"].detach().contiguous()
497
498         assert desc0.shape[-1] == self.conf.input_dim
499         assert desc1.shape[-1] == self.conf.input_dim
500
501         if torch.is_autocast_enabled():
502             desc0 = desc0.half()
503             desc1 = desc1.half()
504
505         mask0, mask1 = None, None
506         c = max(m, n)
507         do_compile = self.static_lengths and c <= max(self.static_lengths)

```

```

508     if do_compile:
509         kn = min([k for k in self.static_lengths if k >= c])
510         desc0, mask0 = pad_to_length(desc0, kn)
511         desc1, mask1 = pad_to_length(desc1, kn)
512         kpts0, _ = pad_to_length(kpts0, kn)
513         kpts1, _ = pad_to_length(kpts1, kn)
514     desc0 = self.input_proj(desc0)
515     desc1 = self.input_proj(desc1)
516     # cache positional embeddings
517     encoding0 = self.posenc(kpts0)
518     encoding1 = self.posenc(kpts1)
519
520     # GNN + final_proj + assignment
521     do_early_stop = self.conf.depth_confidence > 0
522     do_point_pruning = self.conf.width_confidence > 0 and not do_compile
523     pruning_th = self.pruning_min_kpts(device)
524     if do_point_pruning:
525         ind0 = torch.arange(0, m, device=device)[None]
526         ind1 = torch.arange(0, n, device=device)[None]
527         # We store the index of the layer at which pruning is detected.
528         prune0 = torch.ones_like(ind0)
529         prune1 = torch.ones_like(ind1)
530     token0, token1 = None, None
531     for i in range(self.conf.n_layers):
532         if desc0.shape[1] == 0 or desc1.shape[1] == 0: # no keypoints
533             break
534         desc0, desc1 = self.transformers[i](
535             desc0, desc1, encoding0, encoding1, mask0=mask0, mask1=mask1
536         )
537         if i == self.conf.n_layers - 1:
538             continue # no early stopping or adaptive width at last layer
539
540         if do_early_stop:
541             token0, token1 = self.token_confidence[i](desc0, desc1)
542             if self.check_if_stop(token0[...,:m], token1[...,:n], i, m + n):
543                 break
544         if do_point_pruning and desc0.shape[-2] > pruning_th:
545             scores0 = self.log_assignment[i].get_matchability(desc0)
546             prunemask0 = self.get_pruning_mask(token0, scores0, i)
547             keep0 = torch.where(prunemask0)[1]
548             ind0 = ind0.index_select(1, keep0)
549             desc0 = desc0.index_select(1, keep0)
550             encoding0 = encoding0.index_select(-2, keep0)
551             prune0[:, ind0] += 1
552         if do_point_pruning and desc1.shape[-2] > pruning_th:
553             scores1 = self.log_assignment[i].get_matchability(desc1)
554             prunemask1 = self.get_pruning_mask(token1, scores1, i)
555             keep1 = torch.where(prunemask1)[1]
556             ind1 = ind1.index_select(1, keep1)
557             desc1 = desc1.index_select(1, keep1)
558             encoding1 = encoding1.index_select(-2, keep1)
559             prune1[:, ind1] += 1
560
561     if desc0.shape[1] == 0 or desc1.shape[1] == 0: # no keypoints
562         m0 = desc0.new_full((b, m), -1, dtype=torch.long)
563         m1 = desc1.new_full((b, n), -1, dtype=torch.long)
564         mscores0 = desc0.new_zeros((b, m))
565         mscores1 = desc1.new_zeros((b, n))
566         matches = desc0.new_empty((b, 0, 2), dtype=torch.long)
567         mscores = desc0.new_empty((b, 0))
568         if not do_point_pruning:
569             prune0 = torch.ones_like(mscores0) * self.conf.n_layers
570             prune1 = torch.ones_like(mscores1) * self.conf.n_layers
571         return {
572             "matches0": m0,
573             "matches1": m1,
574             "matching_scores0": mscores0,
575             "matching_scores1": mscores1,
576             "stop": i + 1,
577             "matches": matches,
578             "scores": mscores,
579             "prune0": prune0,
580             "prune1": prune1,
581         }
582
583     desc0, desc1 = desc0[..., :m, :], desc1[..., :n, :] # remove padding
584     scores, _ = self.log_assignment[i](desc0, desc1)
585     m0, m1, mscores0, mscores1 = filter_matches(scores, self.conf.filter_threshold)
586     matches, mscores = [], []
587     for k in range(b):
588         valid = m0[k] > -1
589         m_indices_0 = torch.where(valid)[0]
590         m_indices_1 = m0[k][valid]
591         if do_point_pruning:
592             m_indices_0 = ind0[k, m_indices_0]

```



```

593         m_indices_1 = ind1[k, m_indices_1]
594         matches.append(torch.stack([m_indices_0, m_indices_1], -1))
595         mscores.append(mscores0[k][valid])
596
597     # TODO: Remove when hloc switches to the compact format.
598     if do_point_pruning:
599         m0_ = torch.full((b, m), -1, device=m0.device, dtype=m0.dtype)
600         m1_ = torch.full((b, n), -1, device=m1.device, dtype=m1.dtype)
601         m0_[:, ind0] = torch.where(m0 == -1, -1, ind1.gather(1, m0.clamp(min=0)))
602         m1_[:, ind1] = torch.where(m1 == -1, -1, ind0.gather(1, m1.clamp(min=0)))
603         mscores0_ = torch.zeros((b, m), device=mscores0.device)
604         mscores1_ = torch.zeros((b, n), device=mscores1.device)
605         mscores0_[:, ind0] = mscores0
606         mscores1_[:, ind1] = mscores1
607         m0, m1, mscores0, mscores1 = m0_, m1_, mscores0_, mscores1_
608     else:
609         prune0 = torch.ones_like(mscores0) * self.conf.n_layers
610         prune1 = torch.ones_like(mscores1) * self.conf.n_layers
611
612     return {
613         "matches0": m0,
614         "matches1": m1,
615         "matching_scores0": mscores0,
616         "matching_scores1": mscores1,
617         "stop": i + 1,
618         "matches": matches,
619         "scores": mscores,
620         "prune0": prune0,
621         "prune1": prune1,
622     }
623
624 def confidence_threshold(self, layer_index: int) -> float:
625     """scaled confidence threshold"""
626     threshold = 0.8 + 0.1 * np.exp(-4.0 * layer_index / self.conf.n_layers)
627     return np.clip(threshold, 0, 1)
628
629 def get_pruning_mask(
630     self, confidences: torch.Tensor, scores: torch.Tensor, layer_index: int
631 ) -> torch.Tensor:
632     """mask points which should be removed"""
633     keep = scores > (1 - self.conf.width_confidence)
634     if confidences is not None: # Low-confidence points are never pruned.
635         keep |= confidences <= self.confidence_thresholds[layer_index]
636     return keep
637
638 def check_if_stop(
639     self,
640     confidences0: torch.Tensor,
641     confidences1: torch.Tensor,
642     layer_index: int,
643     num_points: int,
644 ) -> torch.Tensor:
645     """evaluate stopping condition"""
646     confidences = torch.cat([confidences0, confidences1], -1)
647     threshold = self.confidence_thresholds[layer_index]
648     ratio_confident = 1.0 - (confidences < threshold).float().sum() / num_points
649     return ratio_confident > self.conf.depth_confidence
650
651 def pruning_min_kpts(self, device: torch.device):
652     if self.conf.flash and FLASH_AVAILABLE and device.type == "cuda":
653         return self.pruning_keypoint_thresholds["flash"]
654     else:
655         return self.pruning_keypoint_thresholds[device.type]

```



```

0 LightGlue/lightglue/sift.py
1 import warnings
2
3 import cv2
4 import numpy as np
5 import torch
6 from kornia.color import rgb_to_grayscale
7 from packaging import version
8
9 try:
10     import pycolmap
11 except ImportError:
12     pycolmap = None
13
14 from .utils import Extractor
15
16
17 def filter_dog_point(points, scales, angles, image_shape, nms_radius, scores=None):
18     h, w = image_shape
19     ij = np.round(points - 0.5).astype(int).T[::-1]
20
21     # Remove duplicate points (identical coordinates).
22     # Pick highest scale or score
23     s = scales if scores is None else scores
24     buffer = np.zeros((h, w))
25     np.maximum.at(buffer, tuple(ij), s)
26     keep = np.where(buffer[tuple(ij)] == s)[0]
27
28     # Pick lowest angle (arbitrary).
29     ij = ij[:, keep]
30     buffer[:] = np.inf
31     o_abs = np.abs(angles[keep])
32     np.minimum.at(buffer, tuple(ij), o_abs)
33     mask = buffer[tuple(ij)] == o_abs
34     ij = ij[:, mask]
35     keep = keep[mask]
36
37     if nms_radius > 0:
38         # Apply NMS on the remaining points
39         buffer[:] = 0
40         buffer[tuple(ij)] = s[keep] # scores or scale
41
42         local_max = torch.nn.functional.max_pool2d(
43             torch.from_numpy(buffer).unsqueeze(0),
44             kernel_size=nms_radius * 2 + 1,
45             stride=1,
46             padding=nms_radius,
47         ).squeeze(0)
48         is_local_max = buffer == local_max.numpy()
49         keep = keep[is_local_max[tuple(ij)]]
50     return keep
51
52
53 def sift_to_rootsift(x: torch.Tensor, eps=1e-6) -> torch.Tensor:
54     x = torch.nn.functional.normalize(x, p=1, dim=-1, eps=eps)
55     x.clip_(min=eps).sqrt_()
56     return torch.nn.functional.normalize(x, p=2, dim=-1, eps=eps)
57
58
59 def run_opencv_sift(features: cv2.Feature2D, image: np.ndarray) -> np.ndarray:
60     """
61     Detect keypoints using OpenCV Detector.
62     Optionally, perform description.
63     Args:
64         features: OpenCV based keypoints detector and descriptor
65         image: Grayscale image of uint8 data type
66     Returns:
67         keypoints: 1D array of detected cv2.KeyPoint
68         scores: 1D array of responses
69         descriptors: 1D array of descriptors
70     """
71     detections, descriptors = features.detectAndCompute(image, None)
72     points = np.array([k.pt for k in detections], dtype=np.float32)
73     scores = np.array([k.response for k in detections], dtype=np.float32)
74     scales = np.array([k.size for k in detections], dtype=np.float32)
75     angles = np.deg2rad(np.array([k.angle for k in detections], dtype=np.float32))
76     return points, scores, scales, angles, descriptors
77
78
79 class SIFT(Extractor):
80     default_conf = {
81         "rootsift": True,
82         "nms_radius": 0, # None to disable filtering entirely.

```

```

83     "max_num_keypoints": 4096,
84     "backend": "opencv", # in {opencv, pycolmap, pycolmap_cpu, pycolmap_cuda}
85     "detection_threshold": 0.0066667, # from COLMAP
86     "edge_threshold": 10,
87     "first_octave": -1, # only used by pycolmap, the default of COLMAP
88     "num_octaves": 4,
89 }
90
91 preprocess_conf = {
92     "resize": 1024,
93 }
94
95 required_data_keys = ["image"]
96
97 def __init__(self, **conf):
98     super().__init__(**conf) # Update with default configuration.
99     backend = self.conf.backend
100     if backend.startswith("pycolmap"):
101         if pycolmap is None:
102             raise ImportError(
103                 "Cannot find module pycolmap: install it with pip"
104                 "or use backend=opencv."
105             )
106         options = {
107             "peak_threshold": self.conf.detection_threshold,
108             "edge_threshold": self.conf.edge_threshold,
109             "first_octave": self.conf.first_octave,
110             "num_octaves": self.conf.num_octaves,
111             "normalization": pycolmap.Normalization.L2, # L1_ROOT is buggy.
112         }
113         device = (
114             "auto" if backend == "pycolmap" else backend.replace("pycolmap_", "")
115         )
116         if (
117             backend == "pycolmap_cpu" or not pycolmap.has_cuda
118         ) and pycolmap.__version__ < "0.5.0":
119             warnings.warn(
120                 "The pycolmap CPU SIFT is buggy in version < 0.5.0, "
121                 "consider upgrading pycolmap or use the CUDA version.",
122                 stacklevel=1,
123             )
124         else:
125             options["max_num_features"] = self.conf.max_num_keypoints
126             self.sift = pycolmap.Sift(options=options, device=device)
127     elif backend == "opencv":
128         self.sift = cv2.SIFT_create(
129             contrastThreshold=self.conf.detection_threshold,
130             nfeatures=self.conf.max_num_keypoints,
131             edgeThreshold=self.conf.edge_threshold,
132             nOctaveLayers=self.conf.num_octaves,
133         )
134     else:
135         backends = {"opencv", "pycolmap", "pycolmap_cpu", "pycolmap_cuda"}
136         raise ValueError(
137             f"Unknown backend: {backend} not in " f"{{{', '.join(backends)}}}."
138         )
139
140 def extract_single_image(self, image: torch.Tensor):
141     image_np = image.cpu().numpy().squeeze(0)
142
143     if self.conf.backend.startswith("pycolmap"):
144         if version.parse(pycolmap.__version__) >= version.parse("0.5.0"):
145             detections, descriptors = self.sift.extract(image_np)
146             scores = None # Scores are not exposed by COLMAP anymore.
147         else:
148             detections, scores, descriptors = self.sift.extract(image_np)
149             keypoints = detections[:, :2] # Keep only (x, y).
150             scales, angles = detections[:, -2:].T
151             if scores is not None and (
152                 self.conf.backend == "pycolmap_cpu" or not pycolmap.has_cuda
153             ):
154                 # Set the scores as a combination of abs. response and scale.
155                 scores = np.abs(scores) * scales
156     elif self.conf.backend == "opencv":
157         # TODO: Check if opencv keypoints are already in corner convention
158         keypoints, scores, scales, angles, descriptors = run_opencv_sift(
159             self.sift, (image_np * 255.0).astype(np.uint8)
160         )
161     pred = {
162         "keypoints": keypoints,
163         "scales": scales,
164         "oris": angles,
165         "descriptors": descriptors,
166     }
167     if scores is not None:

```

```

168         pred["keypoint_scores"] = scores
169
170     # sometimes pycolmap returns points outside the image. We remove them
171     if self.conf.backend.startswith("pycolmap"):
172         is_inside = (
173             pred["keypoints"] + 0.5 < np.array([image_np.shape[-2:][::-1]])
174         ).all(-1)
175         pred = {k: v[is_inside] for k, v in pred.items()}
176
177     if self.conf.nms_radius is not None:
178         keep = filter_dog_point(
179             pred["keypoints"],
180             pred["scales"],
181             pred["oris"],
182             image_np.shape,
183             self.conf.nms_radius,
184             scores=pred.get("keypoint_scores"),
185         )
186         pred = {k: v[keep] for k, v in pred.items()}
187
188     pred = {k: torch.from_numpy(v) for k, v in pred.items()}
189     if scores is not None:
190         # Keep the k keypoints with highest score
191         num_points = self.conf.max_num_keypoints
192         if num_points is not None and len(pred["keypoints"]) > num_points:
193             indices = torch.topk(pred["keypoint_scores"], num_points).indices
194             pred = {k: v[indices] for k, v in pred.items()}
195
196     return pred
197
198 def forward(self, data: dict) -> dict:
199     image = data["image"]
200     if image.shape[1] == 3:
201         image = rgb_to_grayscale(image)
202     device = image.device
203     image = image.cpu()
204     pred = []
205     for k in range(len(image)):
206         img = image[k]
207         if "image_size" in data.keys():
208             # avoid extracting points in padded areas
209             w, h = data["image_size"][k]
210             img = img[:, :h, :w]
211         p = self.extract_single_image(img)
212         pred.append(p)
213     pred = {k: torch.stack([p[k] for p in pred], 0).to(device) for k in pred[0]}
214     if self.conf.rootsift:
215         pred["descriptors"] = sift_to_rootsift(pred["descriptors"])
216     return pred

```

```

0 LightGlue/lightglue/superpoint.py
1 # %BANNER_BEGIN%
2 # -----
3 # %COPYRIGHT_BEGIN%
4 #
5 # Magic Leap, Inc. ("COMPANY") CONFIDENTIAL
6 #
7 # Unpublished Copyright (c) 2020
8 # Magic Leap, Inc., All Rights Reserved.
9 #
10 # NOTICE: All information contained herein is, and remains the property
11 # of COMPANY. The intellectual and technical concepts contained herein
12 # are proprietary to COMPANY and may be covered by U.S. and Foreign
13 # Patents, patents in process, and are protected by trade secret or
14 # copyright law. Dissemination of this information or reproduction of
15 # this material is strictly forbidden unless prior written permission is
16 # obtained from COMPANY. Access to the source code contained herein is
17 # hereby forbidden to anyone except current COMPANY employees, managers
18 # or contractors who have executed Confidentiality and Non-disclosure
19 # agreements explicitly covering such access.
20 #
21 # The copyright notice above does not evidence any actual or intended
22 # publication or disclosure of this source code, which includes
23 # information that is confidential and/or proprietary, and is a trade
24 # secret, of COMPANY. ANY REPRODUCTION, MODIFICATION, DISTRIBUTION,
25 # PUBLIC PERFORMANCE, OR PUBLIC DISPLAY OF OR THROUGH USE OF THIS
26 # SOURCE CODE WITHOUT THE EXPRESS WRITTEN CONSENT OF COMPANY IS
27 # STRICTLY PROHIBITED, AND IN VIOLATION OF APPLICABLE LAWS AND
28 # INTERNATIONAL TREATIES. THE RECEIPT OR POSSESSION OF THIS SOURCE
29 # CODE AND/OR RELATED INFORMATION DOES NOT CONVEY OR IMPLY ANY RIGHTS
30 # TO REPRODUCE, DISCLOSE OR DISTRIBUTE ITS CONTENTS, OR TO MANUFACTURE,
31 # USE, OR SELL ANYTHING THAT IT MAY DESCRIBE, IN WHOLE OR IN PART.
32 #
33 # %COPYRIGHT_END%
34 # -----
35 # %AUTHORS_BEGIN%
36 #
37 # Originating Authors: Paul-Edouard Sarlin
38 #
39 # %AUTHORS_END%
40 # -----*/
41 # %BANNER_END%
42
43 # Adapted by Remi Pautrat, Philipp Lindenberger
44
45 import torch
46 from kornia.color import rgb_to_grayscale
47 from torch import nn
48
49 from .utils import Extractor
50
51
52 def simple_nms(scores, nms_radius: int):
53     """Fast Non-maximum suppression to remove nearby points"""
54     assert nms_radius >= 0
55
56     def max_pool(x):
57         return torch.nn.functional.max_pool2d(
58             x, kernel_size=nms_radius * 2 + 1, stride=1, padding=nms_radius
59         )
60
61     zeros = torch.zeros_like(scores)
62     max_mask = scores == max_pool(scores)
63     for _ in range(2):
64         supp_mask = max_pool(max_mask.float()) > 0
65         supp_scores = torch.where(supp_mask, zeros, scores)
66         new_max_mask = supp_scores == max_pool(supp_scores)
67         max_mask = max_mask | (new_max_mask & (~supp_mask))
68     return torch.where(max_mask, scores, zeros)
69
70
71 def top_k_keypoints(keypoints, scores, k):
72     if k >= len(keypoints):
73         return keypoints, scores
74     scores, indices = torch.topk(scores, k, dim=0, sorted=True)
75     return keypoints[indices], scores
76
77
78 def sample_descriptors(keypoints, descriptors, s: int = 8):
79     """Interpolate descriptors at keypoint locations"""
80     b, c, h, w = descriptors.shape
81     keypoints = keypoints - s / 2 + 0.5
82     keypoints /= torch.tensor(

```

```

83         [(w * s - s / 2 - 0.5), (h * s - s / 2 - 0.5)],
84     ).to(
85         keypoints
86     )[None]
87     keypoints = keypoints * 2 - 1 # normalize to (-1, 1)
88     args = {"align_corners": True} if torch.__version__ >= "1.3" else {}
89     descriptors = torch.nn.functional.grid_sample(
90         descriptors, keypoints.view(b, 1, -1, 2), mode="bilinear", **args
91     )
92     descriptors = torch.nn.functional.normalize(
93         descriptors.reshape(b, c, -1), p=2, dim=1
94     )
95     return descriptors
96
97
98 class SuperPoint(Extractor):
99     """SuperPoint Convolutional Detector and Descriptor
100
101     SuperPoint: Self-Supervised Interest Point Detection and
102     Description. Daniel DeTone, Tomasz Malisiewicz, and Andrew
103     Rabinovich. In CVPRW, 2019. https://arxiv.org/abs/1712.07629
104
105     """
106
107     default_conf = {
108         "descriptor_dim": 256,
109         "nms_radius": 4,
110         "max_num_keypoints": None,
111         "detection_threshold": 0.0005,
112         "remove_borders": 4,
113     }
114
115     preprocess_conf = {
116         "resize": 1024,
117     }
118
119     required_data_keys = ["image"]
120
121     def __init__(self, **conf):
122         super().__init__(**conf) # Update with default configuration.
123         self.relu = nn.ReLU(inplace=True)
124         self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
125         c1, c2, c3, c4, c5 = 64, 64, 128, 128, 256
126
127         self.conv1a = nn.Conv2d(1, c1, kernel_size=3, stride=1, padding=1)
128         self.conv1b = nn.Conv2d(c1, c1, kernel_size=3, stride=1, padding=1)
129         self.conv2a = nn.Conv2d(c1, c2, kernel_size=3, stride=1, padding=1)
130         self.conv2b = nn.Conv2d(c2, c2, kernel_size=3, stride=1, padding=1)
131         self.conv3a = nn.Conv2d(c2, c3, kernel_size=3, stride=1, padding=1)
132         self.conv3b = nn.Conv2d(c3, c3, kernel_size=3, stride=1, padding=1)
133         self.conv4a = nn.Conv2d(c3, c4, kernel_size=3, stride=1, padding=1)
134         self.conv4b = nn.Conv2d(c4, c4, kernel_size=3, stride=1, padding=1)
135
136         self.convPa = nn.Conv2d(c4, c5, kernel_size=3, stride=1, padding=1)
137         self.convPb = nn.Conv2d(c5, 65, kernel_size=1, stride=1, padding=0)
138
139         self.convDa = nn.Conv2d(c4, c5, kernel_size=3, stride=1, padding=1)
140         self.convDb = nn.Conv2d(
141             c5, self.conf.descriptor_dim, kernel_size=1, stride=1, padding=0
142         )
143
144         url = "https://github.com/cvg/LightGlue/releases/download/v0.1_arxiv/superpoint_v1.pth" # noqa
145         self.load_state_dict(torch.hub.load_state_dict_from_url(url))
146
147         if self.conf.max_num_keypoints is not None and self.conf.max_num_keypoints <= 0:
148             raise ValueError("max_num_keypoints must be positive or None")
149
150     def forward(self, data: dict) -> dict:
151         """Compute keypoints, scores, descriptors for image"""
152         for key in self.required_data_keys:
153             assert key in data, f"Missing key {key} in data"
154         image = data["image"]
155         if image.shape[1] == 3:
156             image = rgb_to_grayscale(image)
157
158         # Shared Encoder
159         x = self.relu(self.conv1a(image))
160         x = self.relu(self.conv1b(x))
161         x = self.pool(x)
162         x = self.relu(self.conv2a(x))
163         x = self.relu(self.conv2b(x))
164         x = self.pool(x)
165         x = self.relu(self.conv3a(x))
166         x = self.relu(self.conv3b(x))
167         x = self.pool(x)

```

```

168 x = self.relu(self.conv4a(x))
169 x = self.relu(self.conv4b(x))
170
171 # Compute the dense keypoint scores
172 cPa = self.relu(self.convPa(x))
173 scores = self.convPb(cPa)
174 scores = torch.nn.functional.softmax(scores, 1)[: , :-1]
175 b, _, h, w = scores.shape
176 scores = scores.permute(0, 2, 3, 1).reshape(b, h, w, 8, 8)
177 scores = scores.permute(0, 1, 3, 2, 4).reshape(b, h * 8, w * 8)
178 scores = simple_nms(scores, self.conf.nms_radius)
179
180 # Discard keypoints near the image borders
181 if self.conf.remove_borders:
182     pad = self.conf.remove_borders
183     scores[:, :pad] = -1
184     scores[:, :, :pad] = -1
185     scores[:, -pad:] = -1
186     scores[:, :, -pad:] = -1
187
188 # Extract keypoints
189 best_kp = torch.where(scores > self.conf.detection_threshold)
190 scores = scores[best_kp]
191
192 # Separate into batches
193 keypoints = [
194     torch.stack(best_kp[1:3], dim=-1)[best_kp[0] == i] for i in range(b)
195 ]
196 scores = [scores[best_kp[0] == i] for i in range(b)]
197
198 # Keep the k keypoints with highest score
199 if self.conf.max_num_keypoints is not None:
200     keypoints, scores = list(
201         zip(
202             *[
203                 top_k_keypoints(k, s, self.conf.max_num_keypoints)
204                 for k, s in zip(keypoints, scores)
205             ]
206         )
207     )
208
209 # Convert (h, w) to (x, y)
210 keypoints = [torch.flip(k, [1]).float() for k in keypoints]
211
212 # Compute the dense descriptors
213 cDa = self.relu(self.convDa(x))
214 descriptors = self.convDb(cDa)
215 descriptors = torch.nn.functional.normalize(descriptors, p=2, dim=1)
216
217 # Extract descriptors
218 descriptors = [
219     sample_descriptors(k[None], d[None], 8)[0]
220     for k, d in zip(keypoints, descriptors)
221 ]
222
223 return {
224     "keypoints": torch.stack(keypoints, 0),
225     "keypoint_scores": torch.stack(scores, 0),
226     "descriptors": torch.stack(descriptors, 0).transpose(-1, -2).contiguous(),
227 }

```

```

0 LightGlue/lightglue/utils.py
1 import collections.abc as collections
2 from pathlib import Path
3 from types import SimpleNamespace
4 from typing import Callable, List, Optional, Tuple, Union
5
6 import cv2
7 import kornia
8 import numpy as np
9 import torch
10
11
12 class ImagePreprocessor:
13     default_conf = {
14         "resize": None, # target edge length, None for no resizing
15         "side": "long",
16         "interpolation": "bilinear",
17         "align_corners": None,
18         "antialias": True,
19     }
20
21     def __init__(self, **conf) -> None:
22         super().__init__()
23         self.conf = {**self.default_conf, **conf}
24         self.conf = SimpleNamespace(**self.conf)
25
26     def __call__(self, img: torch.Tensor) -> Tuple[torch.Tensor, torch.Tensor]:
27         """Resize and preprocess an image, return image and resize scale"""
28         h, w = img.shape[-2:]
29         if self.conf.resize is not None:
30             img = kornia.geometry.transform.resize(
31                 img,
32                 self.conf.resize,
33                 side=self.conf.side,
34                 antialias=self.conf.antialias,
35                 align_corners=self.conf.align_corners,
36             )
37         scale = torch.Tensor([img.shape[-1] / w, img.shape[-2] / h]).to(img)
38         return img, scale
39
40
41 def map_tensor(input_, func: Callable):
42     string_classes = (str, bytes)
43     if isinstance(input_, string_classes):
44         return input_
45     elif isinstance(input_, collections.Mapping):
46         return {k: map_tensor(sample, func) for k, sample in input_.items()}
47     elif isinstance(input_, collections.Sequence):
48         return [map_tensor(sample, func) for sample in input_]
49     elif isinstance(input_, torch.Tensor):
50         return func(input_)
51     else:
52         return input_
53
54
55 def batch_to_device(batch: dict, device: str = "cpu", non_blocking: bool = True):
56     """Move batch (dict) to device"""
57
58     def _func(tensor):
59         return tensor.to(device=device, non_blocking=non_blocking).detach()
60
61     return map_tensor(batch, _func)
62
63
64 def rbd(data: dict) -> dict:
65     """Remove batch dimension from elements in data"""
66     return {
67         k: v[0] if isinstance(v, (torch.Tensor, np.ndarray, list)) else v
68         for k, v in data.items()
69     }
70
71
72 def read_image(path: Path, grayscale: bool = False) -> np.ndarray:
73     """Read an image from path as RGB or grayscale"""
74     if not Path(path).exists():
75         raise FileNotFoundError(f"No image at path {path}.")
76     mode = cv2.IMREAD_GRAYSCALE if grayscale else cv2.IMREAD_COLOR
77     image = cv2.imread(str(path), mode)
78     if image is None:
79         raise IOError(f"Could not read image at {path}.")
80     if not grayscale:
81         image = image[..., :-1]
82     return image

```

```

83
84
85 def numpy_image_to_torch(image: np.ndarray) -> torch.Tensor:
86     """Normalize the image tensor and reorder the dimensions."""
87     if image.ndim == 3:
88         image = image.transpose((2, 0, 1)) # HxWxC to CxHxW
89     elif image.ndim == 2:
90         image = image[None] # add channel axis
91     else:
92         raise ValueError(f"Not an image: {image.shape}")
93     return torch.tensor(image / 255.0, dtype=torch.float)
94
95
96 def resize_image(
97     image: np.ndarray,
98     size: Union[List[int], int],
99     fn: str = "max",
100     interp: Optional[str] = "area",
101 ) -> np.ndarray:
102     """Resize an image to a fixed size, or according to max or min edge."""
103     h, w = image.shape[:2]
104
105     fn = {"max": max, "min": min}[fn]
106     if isinstance(size, int):
107         scale = size / fn(h, w)
108         h_new, w_new = int(round(h * scale)), int(round(w * scale))
109         scale = (w_new / w, h_new / h)
110     elif isinstance(size, (tuple, list)):
111         h_new, w_new = size
112         scale = (w_new / w, h_new / h)
113     else:
114         raise ValueError(f"Incorrect new size: {size}")
115     mode = {
116         "linear": cv2.INTER_LINEAR,
117         "cubic": cv2.INTER_CUBIC,
118         "nearest": cv2.INTER_NEAREST,
119         "area": cv2.INTER_AREA,
120     }[interp]
121     return cv2.resize(image, (w_new, h_new), interpolation=mode), scale
122
123
124 def load_image(path: Path, resize: int = None, **kwargs) -> torch.Tensor:
125     image = read_image(path)
126     if resize is not None:
127         image, _ = resize_image(image, resize, **kwargs)
128     return numpy_image_to_torch(image)
129
130
131 class Extractor(torch.nn.Module):
132     def __init__(self, **conf):
133         super().__init__()
134         self.conf = SimpleNamespace(**{**self.default_conf, **conf})
135
136     @torch.no_grad()
137     def extract(self, img: torch.Tensor, **conf) -> dict:
138         """Perform extraction with online resizing"""
139         if img.dim() == 3:
140             img = img[None] # add batch dim
141         assert img.dim() == 4 and img.shape[0] == 1
142         shape = img.shape[-2:][::-1]
143         img, scales = ImagePreprocessor(**{**self.preprocess_conf, **conf})(img)
144         feats = self.forward({"image": img})
145         feats["image_size"] = torch.tensor(shape)[None].to(img).float()
146         feats["keypoints"] = (feats["keypoints"] + 0.5) / scales[None] - 0.5
147         return feats
148
149
150 def match_pair(
151     extractor,
152     matcher,
153     image0: torch.Tensor,
154     image1: torch.Tensor,
155     device: str = "cpu",
156     **preprocess,
157 ):
158     """Match a pair of images (image0, image1) with an extractor and matcher"""
159     feats0 = extractor.extract(image0, **preprocess)
160     feats1 = extractor.extract(image1, **preprocess)
161     matches01 = matcher({"image0": feats0, "image1": feats1})
162     data = [feats0, feats1, matches01]
163     # remove batch dim and move to target device
164     feats0, feats1, matches01 = [batch_to_device(rbd(x), device) for x in data]
165     return feats0, feats1, matches01

```



```

0 LightGlue/lightglue/viz2d.py
1 """
2 2D visualization primitives based on Matplotlib.
3 1) Plot images with `plot_images`.
4 2) Call `plot_keypoints` or `plot_matches` any number of times.
5 3) Optionally: save a .png or .pdf plot (nice in papers!) with `save_plot`.
6 """
7
8 import matplotlib
9 import matplotlib.patheffects as path_effects
10 import matplotlib.pyplot as plt
11 import numpy as np
12 import torch
13
14
15 def cm_RdGn(x):
16     """Custom colormap: red (0) -> yellow (0.5) -> green (1)."""
17     x = np.clip(x, 0, 1)[..., None] * 2
18     c = x * np.array([[0, 1.0, 0]]) + (2 - x) * np.array([[1.0, 0, 0]])
19     return np.clip(c, 0, 1)
20
21
22 def cm_BlRdGn(x_):
23     """Custom colormap: blue (-1) -> red (0.0) -> green (1)."""
24     x = np.clip(x_, 0, 1)[..., None] * 2
25     c = x * np.array([[0, 1.0, 0, 1.0]]) + (2 - x) * np.array([[1.0, 0, 0, 1.0]])
26
27     xn = -np.clip(x_, -1, 0)[..., None] * 2
28     cn = xn * np.array([[0, 0.1, 1, 1.0]]) + (2 - xn) * np.array([[1.0, 0, 0, 1.0]])
29     out = np.clip(np.where(x_[..., None] < 0, cn, c), 0, 1)
30     return out
31
32
33 def cm_prune(x_):
34     """Custom colormap to visualize pruning"""
35     if isinstance(x_, torch.Tensor):
36         x_ = x_.cpu().numpy()
37     max_i = max(x_)
38     norm_x = np.where(x_ == max_i, -1, (x_ - 1) / 9)
39     return cm_BlRdGn(norm_x)
40
41
42 def plot_images(imgs, titles=None, cmaps="gray", dpi=100, pad=0.5, adaptive=True):
43     """Plot a set of images horizontally.
44     Args:
45         imgs: list of NumPy RGB (H, W, 3) or PyTorch RGB (3, H, W) or mono (H, W).
46         titles: a list of strings, as titles for each image.
47         cmaps: colormaps for monochrome images.
48         adaptive: whether the figure size should fit the image aspect ratios.
49     """
50     # conversion to (H, W, 3) for torch.Tensor
51     imgs = [
52         img.permute(1, 2, 0).cpu().numpy()
53         if (isinstance(img, torch.Tensor) and img.dim() == 3)
54         else img
55         for img in imgs
56     ]
57
58     n = len(imgs)
59     if not isinstance(cmaps, (list, tuple)):
60         cmaps = [cmaps] * n
61
62     if adaptive:
63         ratios = [i.shape[1] / i.shape[0] for i in imgs] # W / H
64     else:
65         ratios = [4 / 3] * n
66     figsize = [sum(ratios) * 4.5, 4.5]
67     fig, ax = plt.subplots(
68         1, n, figsize=figsize, dpi=dpi, gridspec_kw={"width_ratios": ratios}
69     )
70     if n == 1:
71         ax = [ax]
72     for i in range(n):
73         ax[i].imshow(imgs[i], cmap=plt.get_cmap(cmaps[i]))
74         ax[i].get_yaxis().set_ticks([])
75         ax[i].get_xaxis().set_ticks([])
76         ax[i].set_axis_off()
77         for spine in ax[i].spines.values(): # remove frame
78             spine.set_visible(False)
79         if titles:
80             ax[i].set_title(titles[i])
81     fig.tight_layout(pad=pad)
82

```

```

83
84 def plot_keypoints(kpts, colors="lime", ps=4, axes=None, a=1.0):
85     """Plot keypoints for existing images.
86     Args:
87         kpts: list of ndarrays of size (N, 2).
88         colors: string, or list of list of tuples (one for each keypoints).
89         ps: size of the keypoints as float.
90     """
91     if not isinstance(colors, list):
92         colors = [colors] * len(kpts)
93     if not isinstance(a, list):
94         a = [a] * len(kpts)
95     if axes is None:
96         axes = plt.gcf().axes
97     for ax, k, c, alpha in zip(axes, kpts, colors, a):
98         if isinstance(k, torch.Tensor):
99             k = k.cpu().numpy()
100         ax.scatter(k[:, 0], k[:, 1], c=c, s=ps, linewidths=0, alpha=alpha)
101
102
103 def plot_matches(kpts0, kpts1, color=None, lw=1.5, ps=4, a=1.0, labels=None, axes=None):
104     """Plot matches for a pair of existing images.
105     Args:
106         kpts0, kpts1: corresponding keypoints of size (N, 2).
107         color: color of each match, string or RGB tuple. Random if not given.
108         lw: width of the lines.
109         ps: size of the end points (no endpoint if ps=0)
110         indices: indices of the images to draw the matches on.
111         a: alpha opacity of the match lines.
112     """
113     fig = plt.gcf()
114     if axes is None:
115         ax = fig.axes
116         ax0, ax1 = ax[0], ax[1]
117     else:
118         ax0, ax1 = axes
119     if isinstance(kpts0, torch.Tensor):
120         kpts0 = kpts0.cpu().numpy()
121     if isinstance(kpts1, torch.Tensor):
122         kpts1 = kpts1.cpu().numpy()
123     assert len(kpts0) == len(kpts1)
124     if color is None:
125         color = matplotlib.cm.hsv(np.random.rand(len(kpts0))).tolist()
126     elif len(color) > 0 and not isinstance(color[0], (tuple, list)):
127         color = [color] * len(kpts0)
128
129     if lw > 0:
130         for i in range(len(kpts0)):
131             line = matplotlib.patches.ConnectionPatch(
132                 xyA=(kpts0[i, 0], kpts0[i, 1]),
133                 xyB=(kpts1[i, 0], kpts1[i, 1]),
134                 coordsA=ax0.transData,
135                 coordsB=ax1.transData,
136                 axesA=ax0,
137                 axesB=ax1,
138                 zorder=1,
139                 color=color[i],
140                 linewidth=lw,
141                 clip_on=True,
142                 alpha=a,
143                 label=None if labels is None else labels[i],
144                 picker=5.0,
145             )
146             line.set_annotation_clip(True)
147             fig.add_artist(line)
148
149     # freeze the axes to prevent the transform to change
150     ax0.autoscale(enable=False)
151     ax1.autoscale(enable=False)
152
153     if ps > 0:
154         ax0.scatter(kpts0[:, 0], kpts0[:, 1], c=color, s=ps)
155         ax1.scatter(kpts1[:, 0], kpts1[:, 1], c=color, s=ps)
156
157
158 def add_text(
159     idx,
160     text,
161     pos=(0.01, 0.99),
162     fs=15,
163     color="w",
164     lcolor="k",
165     lwidth=2,
166     ha="left",
167     va="top",

```

```

168 ):
169     ax = plt.gcf().axes[idx]
170     t = ax.text(
171         *pos, text, fontsize=fs, ha=ha, va=va, color=color, transform=ax.transAxes
172     )
173     if lcolor is not None:
174         t.set_path_effects(
175             [
176                 path_effects.Stroke(linewidth=lwidth, foreground=lcolor),
177                 path_effects.Normal(),
178             ]
179         )
180
181
182 def save_plot(path, **kw):
183     """Save the current figure without any white margin."""
184     plt.savefig(path, bbox_inches="tight", pad_inches=0, **kw)

```

```

0 LightGlue/README.md
1 <p align="center">
2 <h1 align="center"><ins>LightGlue</ins>   
</h1>Local Feature Matching at Light Speed</h1>
3 <p align="center">
4 <a href="https://www.linkedin.com/in/philippilindenberger/">Philipp Lindenberger</a>
5 .
6 <a href="https://psarlin.com/">Paul-Edouard&nbsp;Sarlin</a>
7 .
8 <a href="https://www.microsoft.com/en-us/research/people/mapoll/">Marc&nbsp;Pollefeys</a>
9 </p>
10 <h2 align="center">
11 <p>ICCV 2023</p>
12 <a href="https://arxiv.org/pdf/2306.13643.pdf" align="center">Paper</a> |
13 <a href="https://colab.research.google.com/github/cvg/LightGlue/blob/main/demo.ipynb" align="center">Colab</a> |
14 <a href="https://psarlin.com/assets/LightGlue_ICCV2023_poster_compressed.pdf" align="center">Poster</a> |
15 <a href="https://github.com/cvg/glue-factory" align="center">Train your own!</a>
16 </h2>
17
18 </p>
19 <p align="center">
20 <a href="https://arxiv.org/abs/2306.13643"></a>
21 <br>
22 <em>LightGlue is a deep neural network that matches sparse local features across image pairs.<br>An adaptive mechanism makes it fast for easy pairs (top) and reduces t
23 </p>
24
25 ##
26
27 This repository hosts the inference code of LightGlue, a lightweight feature matcher with high accuracy and blazing fast inference. It takes as input a set of keypoints an
28
29 We release pretrained weights of LightGlue with [SuperPoint] (https://arxiv.org/abs/1712.07629), [DISK] (https://arxiv.org/abs/2006.13566), [ALIKED] (https://arxiv.org/abs/23
30
31 The training and evaluation code can be found in our library [glue-factory] (https://github.com/cvg/glue-factory/).
32
33 ## Installation and demo [!Open In Colab] (https://colab.research.google.com/assets/colab-badge.svg) [https://colab.research.google.com/github/cvg/LightGlue/blob/main/dem
34
35 Install this repo using pip:
36
37 ```bash
38 git clone https://github.com/cvg/LightGlue.git && cd LightGlue
39 python -m pip install -e .
40 ```
41
42 We provide a [demo notebook] (demo.ipynb) which shows how to perform feature extraction and matching on an image pair.
43
44 Here is a minimal script to match two images:
45
46 ```python
47 from lightglue import LightGlue, SuperPoint, DISK, SIFT, ALIKED, DoGHardNet
48 from lightglue.utils import load_image, rbd
49
50 # SuperPoint+LightGlue
51 extractor = SuperPoint(max_num_keypoints=2048).eval().cuda() # load the extractor
52 matcher = LightGlue(features='superpoint').eval().cuda() # load the matcher
53
54 # or DISK+LightGlue, ALIKED+LightGlue or SIFT+LightGlue
55 extractor = DISK(max_num_keypoints=2048).eval().cuda() # load the extractor
56 matcher = LightGlue(features='disk').eval().cuda() # load the matcher
57
58 # load each image as a torch.Tensor on GPU with shape (3,H,W), normalized in [0,1]
59 image0 = load_image('path/to/image_0.jpg').cuda()
60 image1 = load_image('path/to/image_1.jpg').cuda()
61
62 # extract local features
63 feats0 = extractor.extract(image0) # auto-resize the image, disable with resize=None
64 feats1 = extractor.extract(image1)
65
66 # match the features
67 matches01 = matcher({'image0': feats0, 'image1': feats1})
68 feats0, feats1, matches01 = [rbd(x) for x in [feats0, feats1, matches01]] # remove batch dimension
69 matches = matches01['matches'] # indices with shape (K,2)
70 points0 = feats0['keypoints'][matches[:, 0]] # coordinates in image #0, shape (K,2)
71 points1 = feats1['keypoints'][matches[:, 1]] # coordinates in image #1, shape (K,2)
72 ```
73
74 We also provide a convenience method to match a pair of images:
75
76 ```python
77 from lightglue import match_pair
78 feats0, feats1, matches01 = match_pair(extractor, matcher, image0, image1)
79 ```
80
81 ##
82 <p align="center">
83 <a href="https://arxiv.org/abs/2306.13643"></a>
84 <br>
85 <em>LightGlue can adjust its depth (number of layers) and width (number of keypoints) per image pair, with a marginal impact on accuracy.</em>
86 </p>
87
88 ## Advanced configuration
89
90 <details>
91 <summary>[Detail of all parameters - click to expand]</summary>
92
93 - ``n_layers``: Number of stacked self+cross attention layers. Reduce this value for faster inference at the cost of accuracy (continuous red line in the plot above). De
94 - ``flash``: Enable FlashAttention. Significantly increases the speed and reduces the memory consumption without any impact on accuracy. Default: True (LightGlue automat
95 - ``amp``: Enable mixed precision inference. Default: False (off)
96 - ``depth_confidence``: Controls the early stopping. A lower value stops more often at earlier layers. Default: 0.95, disable with -1.
97 - ``width_confidence``: Controls the iterative point pruning. A lower value prunes more points earlier. Default: 0.99, disable with -1.
98 - ``filter_threshold``: Match confidence. Increase this value to obtain less, but stronger matches. Default: 0.1
99
100 </details>
101
102 The default values give a good trade-off between speed and accuracy. To maximize the accuracy, use all keypoints and disable the adaptive mechanisms:
103
104 ```python
105 extractor = SuperPoint(max_num_keypoints=None)
106 matcher = LightGlue(features='superpoint', depth_confidence=-1, width_confidence=-1)
107 ```
108
109 To increase the speed with a small drop of accuracy, decrease the number of keypoints and lower the adaptive thresholds:
110
111 ```python
112 extractor = SuperPoint(max_num_keypoints=1024)
113 matcher = LightGlue(features='superpoint', depth_confidence=0.9, width_confidence=0.95)
114 ```
115
116 The maximum speed is obtained with a combination of:
117 - [FlashAttention] (https://arxiv.org/abs/2205.14135): automatically used when ``torch >= 2.0`` or if [installed from source] (https://github.com/HazyResearch/flash-attent
118 - PyTorch compilation, available when ``torch >= 2.0``:
119
120 ```python
121 matcher = matcher.eval().cuda()
122 matcher.compile(mode='reduce-overhead')
123 ```
124
125 For inputs with fewer than 1536 keypoints (determined experimentally), this compiles LightGlue but disables point pruning (large overhead). For larger input sizes, it autc
126
127 ## Benchmark
128
129 <p align="center">
130 <a></a>
131 <br>
132 <em>Benchmark results on GPU (RTX 3080). With compilation and adaptivity, LightGlue runs at 150 FPS @ 1024 keypoints and 50 FPS @ 4096 keypoints per image. This is a 4-1
133 </p>
134
135 <p align="center">
136 <a></a>
137 <br>
138 <em>Benchmark results on CPU (Intel Xeon W-2175M). With compilation and adaptivity, LightGlue runs at 15 FPS @ 1024 keypoints and 5 FPS @ 4096 keypoints per image. This is a 4-1
139 </p>

```

```

134 <div>
135 <em>Benchmark results on CPU (Intel i7 10700K). LightGlue runs at 20 FPS @ 512 keypoints. </em>
136 </div>
137
138 Obtain the same plots for your setup using our [benchmark script] (benchmark.py):
139 ```
140 python benchmark.py [--device cuda] [--add_superglue] [--num_keypoints 512 1024 2048 4096] [--compile]
141 ```
142
143 <details>
144 <summary>[Performance tip - click to expand]</summary>
145
146 Note: Point pruning introduces an overhead that sometimes outweighs its benefits.
147 Point pruning is thus enabled only when there are more than N keypoints in an image, where N is hardware-dependent.
148 We provide defaults optimized for current hardware (RTX 30xx GPUs).
149 We suggest running the benchmark script and adjusting the thresholds for your hardware by updating LightGlue.pruning_keypoint_thresholds['cuda'].
150
151 </details>
152
153 ## Training and evaluation
154
155 With [Glue Factory] (https://github.com/cvg/glue-factory), you can train LightGlue with your own local features, on your own dataset!
156 You can also evaluate it and other baselines on standard benchmarks like HPatches and MegaDepth.
157
158 ## Other links
159 - [hloc - the visual localization toolbox] (https://github.com/cvg/Hierarchical-Localization/): run LightGlue for Structure-from-Motion and visual localization.
160 - [LightGlue-ONNX] (https://github.com/fabio-sim/LightGlue-ONNX): export LightGlue to the Open Neural Network Exchange (ONNX) format with support for TensorRT and OpenVINO.
161 - [Image Matching WebUI] (https://github.com/Vincentqyw/image-matching-webui): a web GUI to easily compare different matchers, including LightGlue.
162 - [kornia] (https://kornia.readthedocs.io) now exposes LightGlue via the interfaces [LightGlue] (https://kornia.readthedocs.io/en/latest/feature.html#kornia.feature.LightGlue)
163
164 ## BibTeX citation
165 If you use any ideas from the paper or code from this repo, please consider citing:
166
167 ```txt
168 @inproceedings{lindenberger2023lightglue,
169   author    = {Philipp Lindenberger and
170               Paul-Edouard Sarlin and
171               Marc Pollefeys},
172   title     = {{LightGlue: Local Feature Matching at Light Speed}},
173   booktitle = {ICCV},
174   year      = {2023}
175 }
176 ```
177
178 ## License
179 The pre-trained weights of LightGlue and the code provided in this repository are released under the [Apache-2.0 license] (./LICENSE). [DISK] (https://github.com/cvlab-epfl/

```

tree.txt

```

0 LightGlue/
1   ├── assets/
2   │   ├── architecture.svg
3   │   ├── benchmark.png
4   │   ├── benchmark_cpu.png
5   │   ├── DSC_0410.JPG
6   │   ├── DSC_0411.JPG
7   │   ├── easy_hard.jpg
8   │   ├── sacre_coeur1.jpg
9   │   ├── sacre_coeur2.jpg
10  │   └── teaser.svg
11  ├── benchmark.py
12  ├── demo.ipynb
13  ├── demo.py
14  ├── images/
15  │   ├── desk_fix.jpg
16  │   ├── desk_fix2.jpg
17  │   └── desk_vedio.mp4
18  ├── LICENSE
19  ├── lightglue/
20  │   ├── aliked.py
21  │   ├── disk.py
22  │   ├── dog_hardnet.py
23  │   ├── lightglue.py
24  │   ├── sift.py
25  │   ├── superpoint.py
26  │   ├── utils.py
27  │   └── viz2d.py
28  ├── lightglue.egg-info/
29  │   ├── dependency_links.txt
30  │   ├── PKG-INFO
31  │   ├── requires.txt
32  │   ├── SOURCES.txt
33  │   └── top_level.txt
34  ├── pyproject.toml
35  ├── README.md
36  └── requirements.txt

```